

# Composable Statistics

Brian Beckman

May 21, 2019

## Contents

<b>1</b>	<b>COMPOSABLE STATISTICS</b>	<b>2</b>
1.1	CLOJURE . . . . .	2
1.2	<b>TODO</b> HOW TO USE THIS DOCUMENT . . . . .	2
<b>2</b>	<b>INTRODUCTION</b>	<b>2</b>
2.1	TODO: GENERATE NEW RANDOM DATA . . . . .	2
<b>3</b>	<b>RUNNING COUNT</b>	<b>2</b>
3.1	THREAD-SAFE . . . . .	3
3.2	AVOIDING <i>REDUCE</i> . . . . .	3
<b>4</b>	<b>RUNNING MEAN</b>	<b>4</b>
4.1	REMOVING OUTPUT COUPLING . . . . .	4
4.2	NUMERICAL CHECK . . . . .	5
<b>5</b>	<b>CORE.ASYNC</b>	<b>5</b>
5.1	SHALLOW TUTORIAL . . . . .	5
5.2	DEEP TUTORIAL . . . . .	5
<b>6</b>	<b>ASYNC DATA STREAMS</b>	<b>10</b>
6.1	ASYNC RUNNING MEAN . . . . .	11
<b>7</b>	<b>RUNNING STDDEV</b>	<b>13</b>
7.1	BRUTE-FORCE (SCALAR VERSION) . . . . .	13
7.2	DEF Z2S: SMALLER EXAMPLE . . . . .	14
7.3	REALLY DUMB RECURRENCE . . . . .	14
7.4	SCHOOL VARIANCE . . . . .	14
7.5	DEFN MAKE SCHOOL STATS MAPPER . . . . .	15
7.6	DEFN MAKE RECURRENT STATS MAPPER . . . . .	15
7.7	DEFN MAKE WELFORD'S STATS MAPPER . . . . .	16
<b>8</b>	<b>WINDOWED STATISTICS</b>	<b>16</b>
8.1	DEF Z3S: MORE SAMPLE DATA . . . . .	17
8.2	DEFN MAKE SLIDING STATS MAPPER . . . . .	18
<b>9</b>	<b>KALMAN FILTER</b>	<b>19</b>
9.1	BASIC LINEAR ALGEBRA . . . . .	19
9.2	DEFN SYMMETRIC PART . . . . .	19
9.3	DEFN ANTI-SYMMETRIC PART . . . . .	20
9.4	DEFN KALMAN UPDATE: GENERAL EXTENDED KALMAN FILTER . . . . .	22
9.5	DEFN MAKE-KALMAN-MAPPER . . . . .	24

38	<b>10 OZ FOR VISUALIZATION</b>	<b>25</b>
39	10.1 DEFN PLAY DATA . . . . .	25
40	<b>11 GAUSSIAN PROCESSES</b>	<b>26</b>
41	11.1 RECURRENT LINEAR REGRESSION . . . . .	26

## 42 1 COMPOSABLE STATISTICS

43 **Brian Beckman**  
44 **June 2017 - June 2019**

### 45 1.1 CLOJURE

46 We prefer Clojure to Python for this exercise due to Clojure's concurrency primitives, especially atoms and  
47 core.async. Python is growing and improving rapidly, so we may return to it someday.

48 The best sites for learning Clojure by example are [clojuredocs.org](http://clojuredocs.org) and [4clojure.org](http://4clojure.org). A recommended  
49 book is Clojure for the Brave and True.

### 50 1.2 TODO HOW TO USE THIS DOCUMENT

51 This document is adapted from an original Jupyter notebook. I moved to org-mode after a catastrophic  
52 failure of Jupyter. I decided it was more worthwhile to start again from org-mode than from Jupyter.

53 Explain how to run Clojure code inside an org-mode buffer, how to tangle and weave, etc.

54 1. Install leiningen <https://leiningen.org/> (this is all you need for Clojure)

55 Consider doing all Python work in pipenv. It keeps virtual environments outside your local folders. It  
56 works for this `clojupyter` notebook as well.

## 57 2 INTRODUCTION

58 We want to compute descriptive statistics in constant memory. We want exactly the same code to run over  
59 sequences distributed in space as runs over sequences distributed in time. Sequences distributed in space  
60 are vectors, lists, arrays, lazy or not. Sequences distributed over time are asynchronous streams. Descriptive  
61 statistics range from `count`, `mean`, `max`, `min`, and `variance` to Kalman filters and Gaussian processes. We  
62 decouple computation from data delivery by packaging computation in composable functions.

63 Some sample scalar data:

```
64 (def zs [-0.178654, 0.828305, 0.0592247, -0.0121089, -1.48014,  
65         -0.315044, -0.324796, -0.676357, 0.16301, -0.858164])
```

### 66 2.1 TODO: GENERATE NEW RANDOM DATA

## 67 3 RUNNING COUNT

68 The traditional and obvious way with `reduce` and `reductions` ([https://clojuredocs.org/clojure.](https://clojuredocs.org/clojure.core/reduce)  
69 `core/reduce`). *Reduce* takes three arguments: a binary function, an initial value, and a space-sequence of  
70 inputs.

```
71 (reduce  
72   (fn [count datum] (inc count)) ; binary function  
73   0                               ; initial value  
74   zs)                             ; space sequence
```

75 ... with all intermediate results:

```
76 (reductions (fn [c z] (inc c)) 0 zs)
```

### 3.1 THREAD-SAFE

Overkill for sequences in space, but safe for multiple threads from asynchronous streams. It also shows (1) *let-over-lambda* (LOL): closing over mutable state variables, and (2) transactional mutation, i.e., *atomic updates*. LOL is semantically equivalent to data encapsulation in OOP, and transactions are easier to verify than is OOP with locks and mutexes.

The following has a defect: we need `initial-count` both to initialize the `atom` and to initialize the `reduce` call. This defect must be traded off against the generalizable form or *functional type* of the reducible, namely (estimate, measurement) → estimate. We get rid of this defect later.

```
(let [initial-count 0] ; Must use this twice below.
  (reduce
    ; Let-over-lambda (anonymous "object") follows.
    ; "Atom" is a transactional (thread-safe) type in Clojure.
    (let [running-count (atom initial-count)]
      ; That was the "let" of "LOL." Here comes the lambda:
      ; Reducible closure over "running-count."
      (fn [c z] ; Here's the "lambda" of "LOL"
        (swap! running-count inc) ; transactional update
        @running-count))
    initial-count
    zs))
```

Showing all intermediate results:

```
(let [initial-count 0]
  (reductions ; <-- this is the only difference to above
    (let [running-count (atom initial-count)]
      (fn [c z]
        (swap! running-count inc)
        @running-count)) ; ~~> new value for c
    initial-count
    zs))
```

### 3.2 AVOIDING REDUCE

`Reduce` only works in space, not in time. Avoiding `reduce` decouples the statistics code ("business logic") from the space environment ("plumbing"). That space environment delivers data from vectors, lists, etc.). We want to be able to switch out an environment that delivers data from space for an environment that delivers data points `z` from time.

The following is a thread-safe LOL, without `reduce`. We *map* the LOL over a space-sequence in memory to produce exactly the same result as with `reduce`. The mappable LOL does not need an accumulator argument for `count`.

Below, we map *exactly* the same mappable LOL over asynchronous streams.

A subtle defect: the output is still coupled to the computing environment through `print`. We get rid of that, too, below.

```
(dorun ; <-- Discard 'nil's produced by "print."
  (map
    (let [running-count (atom 0)]
      (fn [z] ; <-- one fewer argument
        (swap! running-count inc)
        (print (str @running-count " "))))
    zs))
```

## 4 RUNNING MEAN

Consider the following general scheme for recurrence: *a new statistic is an old statistic plus a correction*. The *correction* is a *gain* times a *residual*. For running mean, the residual is the difference between the new measurement  $z$  and the old mean  $x$ . The gain is  $1/(n+1)$ , where  $n$  is *count-so-far*.  $n$  is a statistic, too, so it is an *old* value, computed and saved before the current observation  $z$  arrived.

/The correction therefore depends only on the new input  $z$  and on old statistics  $x$  and  $n$ . The correction does not depend on new statistics/.

Mathematically, write the general recurrence idea without subscripts as

$$x \leftarrow x + K(z - x)$$

or, with Lamport's notation, wherein new versions of old values get a prime, as an equation

$$x' = x + K(z - x)$$

( $z$  does not have a prime; it is the only exception to the rule that new versions of old quantities have primes).

Contrast the noisy traditional form, which introduces another variable, the index  $n$ . This traditional form is objectively more complicated than either of the two above:

$$x_{n+1} = x_n + K(n)(z_{n+1} - x_n)$$

```
(dorun
  (map
    (let [running-stats (atom {:count 0, :mean 0})]
      (fn [z]
        (let [{x :mean, n :count} @running-stats
              n+1 (inc n) ; cool variable name!
              K (/ 1.0 n+1)]
          (swap! running-stats conj
                 [:count n+1
                  :mean (+ x (* K (- z x)))]))
        (println @running-stats)))
    zs))
```

The swap above calls `conj` on the current contents of the atom `running-stats` and on the rest of the arguments, namely `[:count n+1, :mean ...]`. `conj` is the idiom for “updating” a hashmap, the hashmap in the atom, the hashmap that starts off as `{:count 0, :mean 0}`.

### 4.1 REMOVING OUTPUT COUPLING

Remove `println` from inside the LOL function of  $z$ . Now the LOL function of  $z$  is completely decoupled from its environment. Also, abstract a “factory” method for the LOL, *make-running-stats-mapper*, to clean up the line that does the printing.

#### 4.1.1 MAKE-RUNNING-STATS-MAPPER

```
(defn make-running-stats-mapper []
  (let [running-stats (atom {:count 0 :mean 0 :datum 0})]
    (fn [z]
      (let [{x :mean, n :count, _ :datum} @running-stats
            n+1 (inc n)
            K (/ 1.0 n+1)]
        (swap! running-stats conj
               [:count n+1
                :mean (+ x (* K (- z x)))]))
      [x n]))
```

```

167         [:datum z]))
168     @running-stats)))
169
170 (clojure.pprint/pprint (map (make-running-stats-mapper) zs))

```

## 171 4.2 NUMERICAL CHECK

172 The last value of the running mean is  $-0.279...42$ . Check that against an independent calculation.

### 173 1. DEFN MEAN

```

174 (defn mean [zs] (/ (reduce + zs) (count zs)))
175 (println (mean zs))

```

## 176 5 CORE.ASYNC

177 For data distributed over time, we'll use Clojure's `core.async`. `Core.async` has some subtleties that we  
178 analyze below.

```

179 (require
180   '[clojure.core.async
181     :refer
182     [sliding-buffer dropping-buffer buffer
183      <!!, <!, >!, >!!],
184     go chan onto-chan close!
185     thread alts! alts!! timeout]))

```

### 186 5.1 SHALLOW TUTORIAL

187 <https://github.com/clojure/core.async/blob/master/examples/walkthrough.clj>

### 188 5.2 DEEP TUTORIAL

189 The asynchronous, singleton `go` thread is loaded with very lightweight *pseudothreads* (my terminology, not  
190 standard; most things you will read or see about Clojure.async does not carefully distinguish between  
191 threads and pseudothreads, and I think that's not helpful).

192 Pseudothreads are lightweight state machines that pick up where they left off. It is feasible to have  
193 thousands, even millions of them. Pseudothreads don't block, they *park*. *Parking* and *unparking* are very  
194 fast. We can write clean code with pseudothreads because our code looks like it's blocked waiting for input  
195 or blocked waiting for buffer space. Code with blocking I/O is easy to write and to understand. Code in  
196 `go` forms doesn't actually block, just looks like it.

197 Some details are tricky and definitely not easy to divine from the documentation. Hickey's video from  
198 InfoQ 2013 (<https://www.infoq.com/presentations/core-async-clojure>) is more helpful, but  
199 you can only appreciate the fine points after you've stumbled a bit. I stumbled over the fact that buffered  
200 and unbuffered channels have different synchronization semantics. Syntactically, they look the same, but  
201 you cannot, in general, run the same code over an unbuffered channel that works on a buffered channel.  
202 Hickey says this, but doesn't nail it to the mast; doesn't emphasize it with an example, as I do here in this  
203 deep tutorial. He motivates the entire library with the benefits of first-class queues, but fails to emphasize  
204 that, by default, a channel is not a queue but a blocking rendezvous. He does mention it, but one cannot  
205 fully appreciate the ramifications from a passing glance.

## 5.2.1 COMMUNICATING BETWEEN THREADS AND PSEUDOTHREADS

Write output to unbuffered channel `c` via `>!` on the asynchronous `go` real-thread and read input from the same channel `c` via `<!!` on the UI/REPL `println` real-thread. We'll see later that writing via `>!!` to an unbuffered channel blocks the UI real-thread, so we can't write before reading unbuffered on the UI/REPL real-thread. However, we can write before reading on a non-blocking pseudothread, and no buffer space is needed.

```
(let [c (chan)]          ;; unbuffered chan
      (go (>! c 42))      ;; parks if no space in chan
      (println (<!! c))   ;; blocks UI/REPL until data on c
      (close! c))        ;; idiom; may be harmless overkill
```

In general, single-bang forms work on `go` pseudothreads, and double-bang forms work on real, heavy-weight, Java threads like the UI/REPL thread behind this notebook. In the rest of this notebook, "thread" means "real thread" and we write "pseudothread" explicitly when that's what we mean.

I don't address thread leakage carefully in this tutorial, mostly because I don't yet understand it well. I may overkill by closing channels redundantly.

## 5.2.2 CHANNEL VOODOO FIRST

Writing before reading seems very reasonable, but it does not work on unbuffered channels, as we see below. Before going there, however, let's understand more corners of the example above.

The `go` form itself returns a channel:

```
(clojure.repl/doc go)
```

I believe "the calling thread" above refers to a pseudothread inside the `go` real-thread, but I am not sure because of the ambiguities in the official documentation between "blocking" and "parking" and between "thread" and "well, we don't have a name for them, but Brian calls them 'pseudothreads'."

Is the channel returned by `go` the same channel as `c`?

```
(let [c (chan)]
      (println {:c-channel c})
      (println {:go-channel (go (>! c 42))})
      (println {:c-coughs-up (<!! c)})
      (println {:close-c (close! c)}))
```

No, `c` is a different channel from the one returned by `go`. Consult the documentation for `go` once more:

```
(clojure.repl/doc go)
```

We should be able to read from the channel returned by `go`; call it `d`:

```
(let [c (chan)
      d (go (>! c 42))] ;; 'let' in Clojure is sequential,
                          ;; like 'let*' in Scheme or Common Lisp,
                          ;; so 'd' has a value, here.
      (println {:c-coughs-up (<!! c), ;; won't block
                :d-coughs-up (<!! d)} ;; won't block
      (close! c)
      (close! d))
```

`d`'s coughing up `true` means that the body of the `go`, namely `(>! c 42)` must have returned `true`, because `d` coughs up "the result of the body when completed." Let's see whether our deduction matches documentation for `>!`:

```
(clojure.repl/doc >!)
```

Sure enough. But something important is true and not obvious from this documentation. Writing to `c` inside the `go` block parks the pseudothread because no buffer space is available: `c` was created with a call to `chan` with no arguments, so no buffer space is allocated. Only when reading from `c` does the pseudothread unpark. How? There is no buffer space. Reading on the UI thread manages to short-circuit any need for a buffer and unpark the pseudothread. Such short-circuiting is called a *rendezvous* in the ancient literature of concurrency. Would the pseudothread unpark if we read inside a `go` block and not on the UI thread?

```
(let [c (chan)
      d (go (>! c 42))
      e (go (<! c))]
  (clojure.pprint/pprint {
    :c-channel c, :d-channel d, :e-channel e,
    :e-coughs-up (<!! e), ;; won't block
    :d-coughs-up (<!! d)}) ;; won't block
  (close! c)
  (close! d)
  (close! e))
```

Yes, the pseudothread that parked when 42 is put on `c` via `>!` unparks when 42 is taken off via `<!`. Channel `d` represents the parking step and channel `e` represents the unparking step. All three channels are different.

So now we know how to short-circuit or rendezvous unbuffered channels. In fact, the order of reading and writing (taking and putting) does not matter in the nebulous, asynchronous world of pseudothreads. How Einsteinian is that? The following takes (reads) from `c` on `e` before putting (writing) to `c` on `d`. That's the same as above, only in the opposite order.

```
(let [c (chan)
      e (go (<! c))
      d (go (>! c 42))]
  (clojure.pprint/pprint {
    :c-channel c, :d-channel d, :e-channel e,
    :e-coughs-up (<!! e), ;; won't block
    :d-coughs-up (<!! d)}) ;; won't block
  (close! c)
  (close! d)
  (close! e))
```

### 5.2.3 PUTS BEFORE TAKES CONSIDERED RISKY

`>!!`, by default, blocks if called too early on an unbuffered real thread. We saw above that parked pseudothreads don't block: you can read and write to channels in `go` blocks in any order. However, that's not true with threads that actually block. The documentation is obscure, though not incorrect, about this fact.

```
(clojure.repl/doc >!!)
```

When is “no buffer space available?” It turns out that the default channel constructor makes a channel with no buffer space allocated by default.

```
(clojure.repl/doc chan)
```

We can test the blocking-on-unbuffered case as follows. The following code will block at the line `(>!! c 42)`, as you'll find if you uncomment the code (remove `#_` at the beginning) and run it. You'll have to interrupt the Kernel using the “Kernel” menu at the top of the notebook, and you might have to restart the Kernel, but you should try it once.

```
#_(let [c (chan)]
      (>!! c 42)
      (println (<!! c))
      (close! c))
```

The following variation works fine because we made “buffer space” before writing to the channel. The only difference to the above is the 1 argument to the call of `chan`.

```
(let [c (chan 1)]
  (>!! c 42)
  (println (<!! c))
  (close! c))
```

The difference between the semantics of the prior two examples is not subtle: one hangs the kernel and the other does not. However, the difference in the syntax is subtle and easy to miss.

We can read on the asynchronous `go` pool from the buffered channel `c` because the buffered write (`>!! c`) on the UI thread doesn’t block:

```
(let [c (chan 1)]
  (>!! c 42)
  (println {:go-channel-coughs-up (<!! (go (<! c)))}))
  (close! c))
```

1. ORDER DOESN’T MATTER, SOMEIMES We can do things backwards, reading before writing, even without a buffer. Read from channel (`<! c`) on the `async go` thread “before” writing to (`>!! c 42`) on the REPL / UI thread. “Before,” here, of course, means syntactically or lexically “before,” not temporally.

```
(let [c (chan) ;; NO BUFFER!
      d (go (<! c)) ;; park a pseudothread to read c
      e (>!! c 42)] ;; blocking write un parks c’s pseudothread
  (println {:c-hangs '(<!! c),
            :d-coughs-up (<!! d),
            :what’s-e e})
  (close! c) (close! d))
```

Why did `>!!` produce `true`? Look at docs again:

```
(clojure.repl/doc >!!)
```

Ok, now I fault the documentation. `>!!` will block if there is no buffer space available *and* if there is no *rendezvous* available, that is, no pseudothread parked waiting for `<!`. I have an open question in the Google group for Clojure about this issue with the documentation.

To get the value written in into `c`, we must read `d`. If we tried to read it from `c`, we would block forever because `>!!` blocks when there is no buffer space, and `c` never has buffer space. We get the value out of the `go` nebula by short-circuiting the buffer, by a *rendezvous*, as explained above.

`e`’s being `true` means that `c` wasn’t closed. (`>!! c 42`) should hang.

```
(let [c (chan) ;; NO BUFFER!
      d (go (<! c)) ;; park a pseudothread to read c
      e (>!! c 42) ;; blocking write un parks c’s pseudothread
      f ' (hangs (>!! c 43))] ;; is 'c' closed?
  (println {:c-coughs-up ' (hangs (<!! c)),
            :d-coughs-up (<!! d),
            :what’s-e e,
            :what’s-f f})
  (close! c) (close! d))
```

StackOverflow reveals a way to find out whether a channel is closed by peeking under the covers (<https://stackoverflow.com/questions/24912971>):



```

344 (let [c (chan) ;; NO BUFFER!
345       d (go (<! c)) ;; park a pseudothread to read c
346       e (>!! c 42) ;; blocking write unparks c's pseudothread
347       f (clojure.core.async.impl.protocols/closed? c)]
348   (println {:c-coughs-up ' (hangs (<!! c)),
349            :d-coughs-up (<!! d),
350            :c-is-open-at-e? e,
351            :c-is-open-at-f? f}))
352   (close! c) (close! d))

```

353 2. ORDER DOES MATTER, SOMETIMES Order does matter this time: Writing blocks the UI thread  
 354 without a buffer and no parked read (rendezvous) in the `go` nebula beforehand. I hope you can  
 355 predict that the following will block even before you run it. To be sure, run it, but you'll have to  
 356 interrupt the kernel as before.

```

357 #_(let [c (chan)
358       e (>!! c 42) ;; blocks forever
359       d (go (<! c))]
360   (println {:c-coughs-up ' (this will hang (<!! c)),
361            :d-coughs-up (<!! d),
362            :what's-e e})
363   (close! c) (close! d))

```

#### 364 5.2.4 TIMEOUTS: DON'T BLOCK FOREVER

365 In all cases, blocking calls like `>!!` to unbuffered channels without timeout must appear *last* on the UI,  
 366 non-`go`, thread, and then only if there is some parked pseudothread that's waiting to read the channel by  
 367 short-circuit (rendezvous). If we block too early, we won't get to the line that launches the `async go` nebula  
 368 and parks the short-circuitable pseudothread—parks the rendezvous.

369 The UI thread won't block forever if we add a timeout. `alts!!` is a way to do that. The documenta-  
 370 tion and examples are difficult, but, loosely quoting (emphasis and edits are mine, major ones in square  
 371 brackets):

```

372 (alts!! ports & {:as opts})

```

373 This destructures all keyword options into `opts`. We don't need `opts` or the `:as` keyword below.

374 Completes at most one of several channel operations. [/Not for use inside a (`go ...`) block./]  
 375 **ports is a vector of channel endpoints**, [A channel endpoint is] either a channel to take from  
 376 or a vector of [`channel-to-put-to val-to-put`] pairs, in any combination. Takes will be  
 377 made as if by `<!!`, and puts will be made as if by `>!!`. If more than one port operation is ready,  
 378 a non-deterministic choice will be made unless the `:priority` option is true. If no operation  
 379 is ready and a `:default` value is supplied, [=default-val :default=] will be returned, otherwise  
 380 `alts!!` will [/block/ xxxxpark ?] until the first operation to become ready completes. **Re-**  
 381 **turns [val port] of the completed operation**, where `val` is the value taken for takes, and a  
 382 boolean (`true` unless already closed, as per `put!`) for puts. `opts` are passed as `:key val ...`  
 383 Supported options: `:default val` - the value to use if none of the operations are immediately  
 384 ready `:priority true` - (default `nil`) when `true`, the operations will be tried in order. Note:  
 385 there is no guarantee that the port `exps` or `val` `exps` will be used, nor in what order should they  
 386 be, so they should not be depended upon for side effects.

387 (`alts!! ...`) returns a [`val port`] 2-vector.

388 (`second (alts!! ...)`) is a wrapper of channel `c` We can't write to the resulting `timeout` chan-  
 389 nel because we didn't give it a name.

390 That's a lot of stuff, but we can divine an idiom: pair a channel `c` that *might* block with a fresh `timeout`  
 391 channel in an `alts!!`. At most one will complete. If `c` blocks, the `timeout` will cough up. If `c` coughs up  
 392 before the `timeout` expires, the `timeout` quietly dies (question, is it closed? Will it be left open and leak?)

For a first example, let's make a buffered thread that won't block and pair it with a long timeout. You will see that it's OK to write 43 into this channel (the `[c 43]` term is an implied write; that's clear from the documentation). `c` won't block because it's buffered, it returns immediately, long before the `timeout` could expire.

```
(let [c (chan 1)
      a (alts!! ; outputs a [val port] pair; throw away the val
           ; here are the two channels for `alts!!`
           [[c 43] (timeout 2500)])]
  (clojure.pprint/pprint {:c c, :a a})
  (let [d (go (<! c))]
    (println {:d-returns (<!! d)})
    (close! c))
```

But, if we take away the buffer, the `timeout` channel wins. The only difference to the above is that instead of creating `c` via `(chan 1)`, that is, with a buffer of length 1, we create it with no buffer (and we quoted out the blocking read of `d` with a tick mark).

```
(let [c (chan)
      a (alts!! ; outputs a [val port] pair; throw away the val
           ; here are the two channels for `alts!!`
           [[c 43] (timeout 2500)])]
  (clojure.pprint/pprint {:c c, :a a})
  (let [d (go (<! c))]
    (println {:d-is d})
    ' (println {:d-returns (<!! d)}) ; blocks
    (close! c))
```

## 6 ASYNC DATA STREAMS

The following writes at random times (`>!`) to a parking channel `echo-chan` on an `async go fast` pseudothread. The UI thread block-reads (`<!!`) some data from `echo-chan`. The UI thread leaves values in the channel and thus leaks the channel according to the documentation for `close!` here <https://clojure.github.io/core.async/api-index.html#C>. To prevent the leak permanently, we close the channel explicitly.

```
(def echo-chan (chan))

(doseq [z zs] (go (Thread/sleep (rand 100)) (>! echo-chan z)))
(dotimes [_ 3] (println (<!! echo-chan)))

(println {:echo-chan-closed? (clojure.core.async.impl.protocols/closed? echo-chan)})
(close! echo-chan)
(println {:echo-chan-closed? (clojure.core.async.impl.protocols/closed? echo-chan)})
```

We can chain channels, again with leaks that we explicitly close. Also, we must not `>!` (send) a nil to `repl-chan`, and `<!` can produce nil from `echo-chan` after the timeout and we close `echo-chan`.

```
(clojure.repl/doc <!)
```

Every time you run the block of code below, you will probably get a different result, by design.

```
(def echo-chan (chan))
(def repl-chan (chan))

;; >! chokes on nulls. <! echo-chan can cough up nil if we time out
```

```

439 ;; and close the channel. The following line will throw an exception
440 ;; unless we don't close the channel at the end of this code-block.
441
442 ;; (dotimes [_ 10] (go (>! repl-chan (<! echo-chan))))
443
444 ;; Instead of throwing an exception, just put a random character
445 ;; like \? down the pipe after the echo-chan is closed:
446
447 (dotimes [_ 10] (go (>! repl-chan (or (<! echo-chan) \?))))
448
449 (doseq [z zs] (go (Thread/sleep (rand 100)) (>! echo-chan z)))
450
451 (dotimes [_ 3]
452   (println (<!! (second (alts!! [repl-chan
453                                (timeout 500)])))))
454
455 ;; Alternatively, we can avoid the exception by NOT closing echo-chan.
456 ;; Not closing echo chan will leak it, and that's a lousy idea.
457
458 (close! echo-chan)
459
460 (close! repl-chan)

```

461 Reading from echo-chan may hang the UI thread because the UI thread races the internal go thread  
 462 that reads echo-chan, but the timeout trick works here as above.

```

463 (def echo-chan (chan))
464 (def repl-chan (chan))
465
466 (dotimes [_ 10] (go (>! repl-chan (or (<! echo-chan) \?))))
467 (doseq [z zs] (go (Thread/sleep (rand 100)) (>! echo-chan z)))
468 (dotimes [_ 3]
469   (println (<!! (second (alts!! [echo-chan
470                                (timeout 500)])))))
471
472 (close! echo-chan)
473 (close! repl-chan)

```

474 println on a go pseudoprocess works if we wait long enough. This, of course, is bad practice or "code  
 475 smell."

```

476 (def echo-chan (chan))
477
478 (doseq [z zs] (go (Thread/sleep (rand 100)) (>! echo-chan z)))
479 (dotimes [_ 3] (go (println (<! echo-chan))))
480
481 (Thread/sleep 500) ; no visible output if you remove this line.
482 (close! echo-chan)

```

## 483 6.1 ASYNC RUNNING MEAN

### 484 6.1.1 DEFN ASYNC-RANDOMIZED-SCAN

485 We want running-stats called at random times and with data in random order. A *transducer*, (map  
 486 mapper), lets us collect items off the buffer. The size of the buffer does not matter, but we must specify  
 487 it. Notice that the side-effector effector is passed in, so async-randomized-scan remains decoupled  
 488 from its environment.

In this style of programming, the asynchronous stream might sometimes be called a *functor*, which is anything that's mappable, anything you can map over.

```
(defn async-randomized-scan [zs mapper effector]
  (let [transducer (map mapper)
        ; give buffer length if there is a transducer
        echo-chan (chan (buffer 1) transducer)]
    (doseq [z zs]
      (go (Thread/sleep (rand 100)) (>! echo-chan z)))
    (dotimes [_ (count zs)] (effector (<!! echo-chan))))
  (close! echo-chan)))

(async-randomized-scan zs (make-running-stats-mapper) println)
```

We don't need to explicitly say *buffer*, but I prefer to do.

### 6.1.2 DEFN MAKE SOW REAP

The *effector* above just prints to the console. Suppose we want to save the data?

The following is a version of Wolfram's *Sow* and *Reap* that does not include tags. It uses *atom* for an effectful store because a *let* variable like *result* is not a *var* and *alter-var-root* won't work on *(let [result []] ...)*. An *atom* might be overkill.

*make-sow-reap* returns a message dispatcher in the style of *The Little Schemer*. It responds to namespaced keywords *::sow* and *::reap*. In the case of *::sow*, it returns an *effector* function that *conj*'s its input to the internal result atomically. In the case of *::reap*, it returns the value of the result accumulated so-far.

```
(do (defn make-sow-reap []
      (let [result (atom [])]
        (fn [msg]
          (cond
            (identical? msg ::sow)
            (fn [x] (swap! result #(conj % x)))
            (identical? msg ::reap)
            @result))))))

(let [accumulator (make-sow-reap)]
  (async-randomized-scan zs
    (make-running-stats-mapper)
    (accumulator ::sow))
  (last (accumulator ::reap)))
```

Occasionally, there is some floating-point noise in the very low digits of the mean because *async-randomized-scan* scrambles the order of the inputs. The mean should always be almost equal to  $-0.27947242$ .

### 6.1.3 DEFN ASYNC NON RANDOM SCAN

Of course, the mean of any permutation of the data *zs* is the same, so the order in which data arrive does not change the final result, except for some occasional floating-point noise as mentioned above.

```
(do (defn async-non-random-scan [zs mapper effector]
      (let [transducer (map mapper)
            echo-chan (chan (buffer 1) transducer)]
        (go (doseq [z zs] (>! echo-chan z)))
        (dotimes [_ (count zs)] (effector (<!! echo-chan))))
      (close! echo-chan)))
```

```

536
537 (let [accumulator (make-sow-reap)]
538     (async-non-random-scan zs (make-running-stats-mapper)
539                             (accumulator ::sow))
540     (last (accumulator ::reap))) )

```

#### 541 6.1.4 DEFN SYNC SCAN: WITH TRANSDUCER

542 Here is the modern way, with `transduce`, to reduce over a sequence of data, in order. It's equivalent to  
 543 the non-random async version above. The documentation for `transduce` writes its parameters as `xform f`  
 544 `coll`, and then says

545 reduce with a transformation of `f (xf)`. If `init` is not supplied, `(f)` will be called to produce  
 546 it.

547 Our `xform` is `transducer`, or `(map mapper)`, and our `f` is `conj`, so this is an idiom for mapping  
 548 because `(conj)`, with no arguments, returns `[]`, an appropriate `init`.

```

549 (do (defn sync-scan [zs mapper]
550     (let [transducer (map mapper)]
551         (transduce transducer conj zs)))
552
553     (last (sync-scan zs (make-running-stats-mapper))) )

```

554 We now have complete symmetry between space and time, space represented by the vector `zs` and time  
 555 represented by values on `echo-chan` in random and in non-random order.

## 556 7 RUNNING STDDEV

### 557 7.1 BRUTE-FORCE (SCALAR VERSION)

558 The definition of variance is the following, for  $N > 1$ :

$$\frac{1}{N-1} \sum_{i=1}^N (z_i - \bar{z}_N)^2$$

559 The sum is the *sum of squared residuals*. Each residual is the difference between the  $i$ -th datum  $z_i$  and  
 560 the mean  $\bar{z}_N$  of all  $N$  data in the sample. The outer constant,  $1/(N-1)$  is Bessel's correction.

#### 561 7.1.1 DEFN SSR: SUM OF SQUARED RESIDUALS

562 The following is *brute-force* in the sense that it requires all data up-front so that it can calculate the mean.

```

563 (do (defn ssr [sequ]
564     (let [m (mean sequ)]
565         (reduce #(+ %1 (* (- %2 m) (- %2 m)))
566                 0 sequ)))
567     (ssr zs) )

```

#### 568 7.1.2 DEFN VARIANCE

569 Call `ssr` to compute variance:

```

570 (do
571   (defn variance [sequ]
572     (let [n (count sequ)]
573       (case n
574         0 0
575         1 (first sequ)
576         #_default (/ (ssr sequ) (- n 1.0))))))
577 (variance zs) )

```

## 578 7.2 DEF Z2S: SMALLER EXAMPLE

579 Let's do a smaller example:

```

580 (do (def z2s [55. 89. 144.])
581     (variance z2s) )

```

## 582 7.3 REALLY DUMB RECURRENCE

583 Remember our general form for recurrences,  $x \leftarrow x + K \times (z - x)$ ?

584 We can squeeze running variance into this form in a really dumb way. The following is really dumb  
585 because:

- 586 1. it requires the whole sequence up front, so it doesn't run in constant memory
- 587 2. the intermediate values are meaningless because they refer to the final mean and count, not to the  
588 intermediate ones

589 But, the final value is correct.

```

590 (do (reductions
591     (let [m (mean z2s) ; uh-oh, we refer to _all_ the data ??
592           c (count z2s)]
593       (fn [var z] (+ var (let [r (- z m)] ; residual
594                             (/ (* r r) (- c 1.0))))))
595     0 z2s) )

```

596 That was so dumb that we won't bother with a thread-safe, stateful, or asynchronous form.

## 597 7.4 SCHOOL VARIANCE

598 For an easy, school-level exercise, prove the following equation:

$$\frac{1}{N-1} \sum_{i=1}^N (z_i - \bar{z}_N)^2 = \frac{1}{N-1} \left( \sum_{i=1}^N (z_i^2) - N \bar{z}_N^2 \right)$$

599 Instead of the sum of squared residuals, *ssr*, accumulate the sum of squares, *ssq*.

600 *School variance* is exposed to *catastrophic cancellation* because *ssq* grows quickly. We fix that defect below.

601 We see that something is not best with this form because we don't use the old variance to compute the  
602 new variance. We do better below.

603 Of course, the same mapper works synchronously and asynchronously.

## 7.5 DEFN MAKE SCHOOL STATS MAPPER

and test it both synchronously and asynchronously, randomized and not:

```
(defn make-school-stats-mapper []
  (let [running-stats (atom {:count 0, :mean 0,
                             :variance 0, :ssq 0})]
    (fn [z]
      (let [{x :mean, n :count, s :ssq} @running-stats
            n+1 (inc n)
            K (/ 1.0 n+1)
            r (- z x)
            x' (+ x (* K r)) ;; Isn't it nice we can use prime notation?
            s' (+ s (* z z))]
        (swap! running-stats conj
                [:count n+1]
                [:mean x']
                [:ssq s']
                [:variance (/ (- s' (* n+1 x' x')) (max 1 n))]))
      @running-stats)))

(clojure.pprint/pprint (sync-scan z2s (make-school-stats-mapper)))

(async-randomized-scan z2s (make-school-stats-mapper) println)

(async-non-random-scan z2s (make-school-stats-mapper) println)
```

## 7.6 DEFN MAKE RECURRENT STATS MAPPER

We already know the recurrence for the mean:

$$x \leftarrow x + K \cdot (z - x) = x + \frac{1}{n+1}(z - x)$$

We want a recurrence with a similar form for the variance. It takes a little work to prove, but it's still a school-level exercise.  $K$  remains  $1/(n+1)$ , the value needed for the new mean. We could define a pair of gains, one for the mean and one for the variance, but it would be less pretty.

$$v \leftarrow \frac{(n-1)v + K n (z - x)^2}{\max(1, n)}$$

```
(defn make-recurrent-stats-mapper []
  (let [running-stats (atom {:count 0, :mean 0,
                             :variance 0})]
    (fn [z]
      (let [{x :mean, n :count, v :variance} @running-stats
            n+1 (inc n)
            K (/ 1.0 (inc n))
            r (- z x)
            x' (+ x (* K r))
            ssr (+ (* (- n 1) v) ; old ssr is (* (- n 1) v)
                  (* K n r r))]
        (swap! running-stats conj
                [:count n+1]
                [:mean x']
                [:variance (/ ssr (max 1 n))]))
      @running-stats)))
```

```

649
650 (async-non-random-scan z2s (make-recurrent-stats-mapper) println)

```

## 651 7.7 DEFN MAKE WELFORD'S STATS MAPPER

652 The above is equivalent, algebraically and numerically, to Welford's famous recurrence for the sum of  
 653 squared residuals  $S$ . In recurrences, we want everything on the right-hand sides of equations or left arrows  
 654 to be old, *prior* statistics, except for the new observation / measurement / input  $z$ . Welford's requires  
 655 the new, *posterior* mean on the right-hand side, so it's not as elegant as our recurrence above. However, it is  
 656 easier to remember!

$$S \leftarrow S + (z - x_N)(z - x_{N+1}) = S + (z - x)(z - (x + K(z - x)))$$

```

657 (do (defn make-welfords-stats-mapper []
658       (let [running-stats (atom {:count 0, :mean 0, :variance 0})]
659         (fn [z]
660           (let [{x :mean, n :count, v :variance} @running-stats
661                 n+1 (inc n)
662                 K (/ 1.0 n+1)
663                 r (- z x)
664                 x' (+ x (* K r))
665                 ssr (+ (* (- n 1) v)
666                       ;; only difference to recurrent variance:
667                       (* (- z x) (- z x')))]
668             (swap! running-stats conj
669                   [:count n+1]
670                   [:mean x']
671                   [:variance (/ ssr (max 1 n))]))
672             @running-stats)))
673
674 (async-non-random-scan
675   z2s (make-welfords-stats-mapper) println)

```

## 676 8 WINDOWED STATISTICS

677 Suppose we want running statistics over a history of fixed, finite length. For example, suppose we have  
 678  $N = 10$  data and we want the statistics in a window of length  $w = 3$  behind the current value, inclusively.  
 679 When the first datum arrives, the window and the total include one datum. The window overhangs the  
 680 left until the third datum. When the fourth datum arrives, the window contains three data and the total  
 681 contains four data. After the tenth datum, we may consider three more steps marching the window "off  
 682 the cliff" to the right. The following figure illustrates (the first row corresponds to  $n = 0$ , not to  $n = 1$ ):

683 We won't derive the following formulas, but rather say that they have been vetted at least twice inde-  
 684 pendently (in a C program and in a Mathematica program). The following table shows a unit test that we  
 685 reproduce. The notation is explained after the table.

686 Denote prior statistics by plain variables like  $m$  and corresponding posteriors by the same variables  
 687 with primes like  $m'$ . The posteriors  $j$  and  $u$  do not have a prime.



variable	description
$n$	prior count of data points; equals 0 when considering the first point
$z$	current data point
$w$	fixed, constant, maximum width of window; $w \geq 1$
$j$	posterior number of points left of the window; $j \geq 0$
$u$	posterior number of points including $z$ in the running window; $1 \leq u \leq w$
$m$	prior mean of all points, not including $z$
$m'$	posterior mean of all points including $z$
$m_j$	prior mean of points left of the window, lagging $w$ behind $m$
$m'_j$	posterior mean of points left of the window
$m'_w$	posterior mean of points in the window, including the current point $z$
$v$	prior variance, not including $z$
$v'$	posterior variance of all points including $z$
$v_j$	prior variance of points left of the window, lagging $w$ behind $u_n$
$v'_j$	posterior variance of points left of the window
$v'_w$	posterior variance of points within the window

The recurrences for  $m$ ,  $v$ ,  $m_j$ , and  $v_j$  have only priors (no primes) on their right-hand sides. The values of  $m_w$  and  $v_w$  are not recurrences because the non-primed versions do not appear on the right-hand sides of equations 10 and 13. Those equations are simply transformations of the posteriors (values with primes)  $m'$ ,  $m'_j$ ,  $v'$ , and  $v'_j$ .

\$\$

$$j = \max(0, n + 1 - w) \quad (1)$$

$$u = n - j + 1 \quad (2)$$

$$m' = m + \frac{z - m}{n + 1} \quad (3)$$

$$m'_j = \begin{cases} m_j + \frac{z_j - m_j}{j} & j > 0 \\ 0 & \text{otherwise} \end{cases} \quad (4)$$

$$m'_w = \frac{(n + 1) m' - j m'_j}{u} \quad (5)$$

$$v' = \frac{(n - 1) v + \frac{n}{n+1} (z - m)^2}{\max(1, n)} \quad (6)$$

$$v'_j = \begin{cases} \frac{j-2}{j-1} v_j + \frac{1}{j} (z_j - m_j)^2 & j > 1 \\ 0 & \text{otherwise} \end{cases} \quad (7)$$

$$v'_w = \frac{n v' + (n - w) v'_j + (n + 1) m'^2 - j m_j'^2 - u m_w'^2}{\max(1, u - 1)} \quad (8)$$

\$\$

Here is sample data we can compare with the unit test above.

## 8.1 DEF Z3S: MORE SAMPLE DATA

```
(def z3s [0.857454, 0.312454, 0.705325, 0.839363, 1.63781, 0.699257, -0.340016, -0.213596,
```

The best algorithm we have found for tracking historical data is to keep a FIFO queue in a Clojure *vector* of length  $w$ . This is still constant memory because it depends only on the length  $w$  of the window, not on the length of the data stream.

### 8.1.1 DEFN PUSH TO BACK

```
(defn push-to-back [item vek]
  (conj (vec (drop 1 vek)) item))
```

## 8.2 DEFN MAKE SLIDING STATS MAPPER

```

703
704 (defn make-sliding-stats-mapper [w]
705   (let [running-stats (atom {:n 0, :m 0, :v 0,
706                               :win (vec (repeat w 0)),
707                               :mw 0, :vw 0,
708                               :mj 0, :vj 0})]
709     (fn [z]
710       (let [{:keys [m n v win mj vj]} @running-stats
711             zj (first win)
712             win' (push-to-back z win)
713             n+1 (double (inc n))
714             n-1 (double (dec n))
715             K (/ 1.0 n+1)
716             Kv (* n K)
717             r (- z m)
718             j (max 0, (- n+1 w))
719             u (- n+1 j)
720             m' (+ m (* K r))
721             rj (- zj mj)
722             mj' (if (> j 0), (+ mj (/ rj j)), 0)
723             mw' (/ (- (* n+1 m') (* j mj')) u)
724             v' (/ (+ (* n-1 v) (* Kv r r))
725                   (max 1 n))
726             vj' (if (> j 1)
727                   (let [j21 (/ (- j 2.0)
728                                (- j 1.0))]
729                     (+ (* j21 vj)
730                       (/ (* rj rj) j)))
731                   0)
732             vw' (let [t1 (- (* n v')
733                             (* (- n w) vj'))
734                      t2 (- (* n+1 m' m')
735                             (* j mj' mj'))
736                      t3 (- (* u mw' mw'))]
737                   (/ (+ t1 t2 t3)
738                     (max 1 (- u 1))))
739             ]
740       (swap! running-stats conj
741             [:n n+1 ]
742             [:m m' ]
743             [:v v' ]
744             [:mj mj' ]
745             [:vj vj' ]
746             [:mw mw' ]
747             [:vw vw' ]
748             [:win win'])
749       @running-stats)))
750
751 (clojure.pprint/print-table
752   [:n :mw :vw]
753   (sync-scan z3s (make-sliding-stats-mapper 3)))

```

... passing the unit test.

## 9 KALMAN FILTER

### 9.1 BASIC LINEAR ALGEBRA

Go for high performance with CUDA or Intel KML later.

Add the following lines to `project.clj` in the directory that contains this org file:

#### 9.1.1 TODO: FULLY LITERATE: TANGLE PROJECT.CLJ

```
[net.mikera/core.matrix "0.62.0"]
[net.mikera/vectorz-clj "0.48.0"]
[org.clojure/algo.generic "0.1.2"]
```

Smoke test:

```
(require '[clojure.core.matrix :as ccm])
(ccm/set-current-implementation :vectorz)
```

```
(ccm/shape
  (ccm/array [[1 2 3]
              [1 3 8]
              [2 7 4]]))
```

Bits and pieces we will need:

```
(ccm/transpose
  (ccm/array [[1 2 3]
              [1 3 8]
              [2 7 4]]))
```

`mmul` is multiadic (takes more than two arguments). This is possible because matrix multiplication is associative.

```
(let [A (ccm/array [[1 2 3]
                    [1 3 8]
                    [2 7 4]])]
  (ccm/mmul (ccm/transpose A) A (ccm/inverse A)))
```

#### 9.1.2 DEFN Linspace

```
(defn linspace
  "A sequence of $n$ equally spaced points in the doubly closed
  interval $[a,b]$, that is, inclusive of both ends."
  [a b n]
  (let [d (/ (- b a) (dec n))]
    (map (fn [x] (+ a (* x d))) (range n))))

(clojure.pprint/pprint (linspace 2 3. 3))
```

### 9.2 DEFN SYMMETRIC PART

```
(do (defn symmetric-part [M]
      (ccm/div (ccm/add M (ccm/transpose M)) 2.0))
    (symmetric-part [[1 2 3]
                     [1 3 8]
                     [2 7 4]]))
```

### 9.3 DEFN ANTI-SYMMETRIC PART

```
(do (defn anti-symmetric-part [M]
      (ccm/div (ccm/sub M (ccm/transpose M)) 2.0))
    (anti-symmetric-part [[1 2 3]
                          [1 3 8]
                          [2 7 4]]))

(let [M [[1 2 3]
         [1 3 8]
         [2 7 4]]]
      (ccm/sub (ccm/add (symmetric-part M)
                        (anti-symmetric-part M))
                M))
```

#### 9.3.1 DEFN MATRIX ALMOST =

```
(require '[clojure.algo.generic.math-functions :as gmf])
```

The following isn't the best solution: neither relative nor absolute differences are robust. Units in Last Place (ULP) are a better criterion, however, this will unblock us for now.

```
(do (defn matrix-almost=
      ([m1 m2 eps]
       "Checks for near equality against a given absolute difference."
       (mapv (fn [row1 row2]
                (mapv (fn [e1 e2] (gmf/approx= e1 e2 eps))
                      row1 row2))
              m1 m2))
      ([m1 m2]
       "Checks for near equality against a default absolute difference of 1.0e-9"
       (matrix-almost= m1 m2 1.0e-9)))

(let [M [[1 2 3]
         [1 3 8]
         [2 7 4]]]
      (matrix-almost= (ccm/add (symmetric-part M)
                              (anti-symmetric-part M))
                      M))
```

#### 9.3.2 DEFN SIMILARITY TRANSFORM

```
(defn similarity-transform [A M]
  (ccm/mmul A M (ccm/transpose A)))
```

#### 9.3.3 VECTORS, ROW VECTORS, COLUMN VECTORS

The library (like many others) is loose about matrices times vectors.

```
(ccm/mmul
  (ccm/matrix [[1 2 3]
               [1 3 8]
               [2 7 4]])
  (ccm/array [22 23 42]))
```

Pedantically, a matrix should only be allowed to left-multiply a column vector, i.e., a  $1 \times 3$  matrix. The Clojure library handles this case.

```

840 (ccm/mmul
841   (ccm/matrix [[1 2 3]
842               [1 3 8]
843               [2 7 4]]))
844   (ccm/array [[22] [23] [42]]))

```

845 Non-pedantic multiplication of a vector on the right by a matrix:

```

846 (ccm/mmul
847   (ccm/array [22 23 42])
848   (ccm/matrix [[1 2 3]
849               [1 3 8]
850               [2 7 4]]))

```

851 Pedantic multiplication of a row vector on the right by a matrix:

```

852 (ccm/mmul
853   (ccm/array [[22 23 42]])
854   (ccm/matrix [[1 2 3]
855               [1 3 8]
856               [2 7 4]]))

```

### 857 9.3.4 SOLVING INSTEAD OF INVERTING

858 Textbooks will tell you that, if you have  $\mathbf{Ax} = \mathbf{b}$  and you want  $\mathbf{x}$ , you should compute  $\mathbf{A}^{-1}\mathbf{b}$ . Don't do this;  
 859 the inverse is numerically risky and almost never needed:

```

860 (ccm/mmul
861   (ccm/inverse
862     (ccm/array [[1 2 3]
863                 [1 3 8]
864                 [2 7 4]]))
865   (ccm/array [22 23 42]))

```

866 Instead, use a linear solver. Almost everywhere that you see  $\mathbf{A}^{-1}\mathbf{b}$ , visualize `solve(A, b)`. You will get  
 867 a more stable answer. Notice the difference in the low-significance digits below. The following is a more  
 868 reliable answer:

```

869 (require '[clojure.core.matrix.linear :as ccml])

870 (ccml/solve
871   (ccm/array [[1 2 3]
872               [1 3 8]
873               [2 7 4]]))
874   (ccm/array [22 23 42]))

875 (ccml/solve
876   (ccm/matrix [[1 2 3]
877               [1 3 8]
878               [2 7 4]]))
879   (ccm/matrix [22 23 42]))

880 (ccm/shape (ccm/matrix [[22] [23] [42]]))

```

### 9.3.5 DEFN SOLVE MATRIX

We need solve to work on matrices:

```
(defn solve-matrix
  "The 'solve' routine in clojure.core.matrix only works on Matrix times Vector.
  We need it to work on Matrix times Matrix. The equation to solve is
  Ann * Xnm = Bnm
  Think of the right-hand side matrix Bnm as a sequence of columns. Iterate over
  its transpose, treating each column as a row, then converting that row to a
  vector, to get the transpose of the solution X."
  [Ann Bnm]
  (ccm/transpose (mapv (partial ccm/solve Ann) (ccm/transpose Bnm))))

(solve-matrix
  (ccm/matrix [[1 2 3]
               [1 3 8]
               [2 7 4]]))

(solve-matrix
  (ccm/matrix [[1 2 3]
               [1 3 8]
               [2 7 4]]))
  (ccm/matrix [[22 44]
               [23 46]
               [42 84]]))
```

### 9.4 DEFN KALMAN UPDATE: GENERAL EXTENDED KALMAN FILTER

Use Clojure's destructuring to write the Kalman filter as a binary function. See <http://vixra.org/abs/1606.0348>

$\mathbf{x}_{n,1}$  denotes a vector  $\mathbf{x}$  with dimension  $n \times 1$ , that is, a column vector of height  $n$ .  $\mathbf{P}_{n,n}$  denotes a covariance matrix of dimension  $n \times n$ , and so on.

The math is as follows (notice step 6 has the same form as all earlier statistics calculations in this document):

Letting inputs:

- $\mathbf{x}_{n,1}$  be the current, best estimate of the  $n$ -dimensional state of a system
- $\mathbf{P}_{n,n}$  be the current, best estimate of the  $n \times n$  covariance of state  $\mathbf{x}_{n,1}$
- $\mathbf{z}_{m,1}$  be the current,  $m$ -dimensional observation
- $\mathbf{H}_{m,n}$  be linearized observation model to be inverted:  $\mathbf{z}_{m,1} = \mathbf{H}_{m,n} \cdot \mathbf{x}_{n,1}$
- $\mathbf{A}_{n,n}$  be linearized dynamics
- $\mathbf{Q}_{n,n}$  be process noise (covariance) accounting for uncertainty in  $\mathbf{A}_{n,n}$
- $\mathbf{R}_{m,m}$  be observation noise (covariance) accounting for uncertainty in  $\mathbf{z}_{m,1}$

and intermediates and outputs:

- $\mathbf{x}'_{n,1}$  (intermediate; *update*) be the estimate of the state after enduring one time step of linearized dynamics

- $\mathbf{x}_{n,1}''$  (output; *prediction*) be the estimate of the state after dynamics and after information from the observation  $\mathbf{z}_{m,1}$
- $\mathbf{P}_{n,n}'$  (intermediate; *update*) be the current, best estimate of the  $n \times n$  covariance of state  $\mathbf{x}_{n,1}$  after dynamics
- $\mathbf{P}_{n,n}''$  (output; *prediction*) be the current, best estimate of the  $n \times n$  covariance of state  $\mathbf{x}_{n,1}$  after dynamics and observation  $\mathbf{z}_{m,1}$

The steps are:

1. *Update state estimate*:  $\mathbf{x}_{n,1}' = \mathbf{A}_{n,n} \mathbf{x}_{n,1}$
2. *Update state covariance*:  $\mathbf{P}_{n,n}' = \mathbf{Q}_{n,n} + (\mathbf{A}_{n,n} \mathbf{P}_{n,n} \mathbf{A}_{n,n}^T)$
3. *Covariance-update scaling matrix*:  $\mathbf{D}_{m,m} = \mathbf{R}_{m,m} + (\mathbf{H}_{m,n} \mathbf{P}_{n,n}' \mathbf{H}_{m,n}^T)$
4. *Kalman gain*:  $\mathbf{K}_{n,m} = \mathbf{P}_{n,n}' \mathbf{H}_{m,n}^T \mathbf{D}_{m,m}^{-1}$   
(a) written as  $\mathbf{K}_{n,m}^T = \text{solve}(\mathbf{D}_{m,m}^T, \mathbf{H}_{m,n} \mathbf{P}_{n,n}')$
5. *Innovation: predicted observation residual*:  $\mathbf{r}_{m,1} = \mathbf{z}_{m,1} - \mathbf{H}_{m,n} \mathbf{x}_{n,1}'$
6. *State prediction*:  $\mathbf{x}_{n,1}'' = \mathbf{x}_{n,1}' + \mathbf{K}_{n,m} \mathbf{r}_{m,1}$
7. *Covariance reduction matrix*:  $\mathbf{L}_{n,n} = \mathbf{I}_{n,n} - \mathbf{K}_{n,m} \mathbf{H}_{m,n}$
8. *Covariance prediction*:  $\mathbf{P}_{n,n}'' = \mathbf{L}_{n,n} \mathbf{P}_{n,n}'$

```
(defn kalman-update [{:keys [xn1 Pnn]} {:keys [zm1 Hmn Ann Qnn Rmm]})
  (let [x' n1 (ccm/mmul Ann xn1) ; Predict state
        P' nn (ccm/add
                Qnn (similarity-transform Ann Pnn)) ; Predict covariance
        Dmm (ccm/add
              Rmm (similarity-transform Hmn P' nn)) ; Gain precursor
        DTmm (ccm/transpose Dmm) ; Support for "solve"
        HP' Tmn (ccm/mmul Hmn (ccm/transpose P' nn)) ; Support for "solve"
        KTmn (solve-matrix DTmm HP' Tmn) ; Eqn 3 of http://vixra.org/abs/1606.0
        Knm (ccm/transpose KTmn) ; Kalman gain
        rml (ccm/sub zm1 (ccm/mmul Hmn x' n1)) ; innovation = predicted obn residual
        x'' n1 (ccm/add x' n1 (ccm/mmul Knm rml)) ; final corrected estimate
        n (ccm/dimension-count xn1 0)
        Lnn (ccm/sub (ccm/identity-matrix n) ; Support for new covariance ...
                     (ccm/mmul Knm Hmn)) ; ... ? catastrophic cancellation ?
        P'' nn (ccm/mmul Lnn P' nn)] ; New covariance
    {:xn1 x'' n1, :Pnn P'' nn}))
```

#### 9.4.1 UNIT TEST

Let the measurement model be a cubic:

```
(defn Hmn-t [t]
  (ccm/matrix [( * t t t) (* t t) t 1 ]))
```

Ground truth state, constant with time in this unit test:

```
(def true-x
  (ccm/array [-5 -4 9 -3]))
```

```

965 (require '[clojure.core.matrix.random :as ccmr])

966 (defn fake [n]
967   (let [times      (range -2.0 2.0 (/ 2.0 n))
968         Hmns       (mapv Hmn-t times)
969         true-zs     (mapv #(ccm/mmul % true-x) Hmns)
970         zmls        (mapv #(ccm/add
971                             % (ccm/array
972                                 [(ccmr/rand-gaussian)]))
973                             true-zs)]
974     {:times times, :Hmns Hmns, :true-zs true-zs, :zmls zmls}))

975 (def test-data (fake 7))

976   A state cluster is a vector of  $\mathbf{x}$  and  $\mathbf{P}$ :

977 (def state-cluster-prior
978   {:xn1 (ccm/array [[0.0] [0.0] [0.0] [0.0]])
979    :Pnn (ccm/mul 1000.0 (ccm/identity-matrix 4))})

980   An obn-cluster is a vector of  $\mathbf{z}$ ,  $\mathbf{H}$ ,  $\mathbf{A}$ ,  $\mathbf{Q}$ , and  $\mathbf{R}$ . Obn is short for observation.

981 (def obn-clusters
982   (let [c (count (:times test-data))]
983     (mapv (fn [zml Hmn Ann Qnn Rmm]
984             {:zml zml, :Hmn Hmn, :Ann Ann, :Qnn Qnn, :Rmm Rmm})
985           (:zmls test-data)
986           (:Hmns test-data)
987           (repeat c (ccm/identity-matrix 4))
988           (repeat c (ccm/zero-matrix 4 4))
989           (repeat c (ccm/identity-matrix 1))
990           )))

991 (clojure.pprint/pprint (reduce kalman-update state-cluster-prior obn-clusters))

992   Notice how close the estimate  $\mathbf{x}_{n \times 1}$  is to the ground truth,  $[-5, -4, 9, -3]$  for  $\mathbf{x}$ . A chi-squared test would
993   be appropriate to complete the verification (TODO).

```

## 994 9.5 DEFN MAKE-KALMAN-MAPPER

995 Just as we did before, we can convert a *foldable* into a *mappable* transducer and bang on an asynchronous  
996 stream of data. This only needs error handling to be deployable at scale. Not to minimize error handling:  
997 it's a big but separable engineering task.

```

998 (do (defn make-kalman-mapper [{:keys [xn1 Pnn]}]
999     ;; let-over-lambda (LOL); here are the Bayesian priors
1000     (let [estimate-and-covariance (atom {:xn1 xn1, ;; prior-estimate
1001                                           :Pnn Pnn, ;; prior-covariance
1002                                           })]
1003       ;; here is the mapper (mappable)
1004       (fn [{:keys [zml Hmn Ann Qnn Rmm]}]
1005         (let [{:keys [xn1 :xn1, Pnn :Pnn]} @estimate-and-covariance]
1006           (let [;; out-dented so we don't go crazy reading it
1007                 x' n1 (ccm/mmul Ann xn1) ; Predict state
1008                 P' nn (ccm/add
1009                       Qnn (similarity-transform Ann Pnn)) ; Predict covariance
1010                 Dmm (ccm/add

```



```

1011      Rmm (similarity-transform Hmn P'nn)) ; Gain precursor
1012      DTmm (ccm/transpose Dmm) ; Support for "solve"
1013      HP'Tmn (ccm/mmul Hmn (ccm/transpose P'nn)) ; Support for "solve"
1014      KTmn (solve-matrix DTmm HP'Tmn) ; Eqn 3 of http://vixra.org/abs/1606.0
1015      Knm (ccm/transpose KTmn) ; Kalman gain
1016      rml (ccm/sub zml (ccm/mmul Hmn x'n1)) ; innovation = predicted obn residual
1017      x''n1 (ccm/add x'n1 (ccm/mmul Knm rml)) ; final corrected estimate
1018      n (ccm/dimension-count xn1 0)
1019      Lnn (ccm/sub (ccm/identity-matrix n) ; Support for new covariance ...
1020              (ccm/mmul Knm Hmn)) ; ... ? catastrophic cancellation ?
1021      P''nn (ccm/mmul Lnn P'nn)]
1022      (swap! estimate-and-covariance conj
1023              [:xn1 x''n1]
1024              [:Pnn P''nn]) ) )
1025      @estimate-and-covariance) ))
1026
1027  ;; The following line maps over a fixed sequence in memory
1028  #_(clojure.pprint/pprint (last
1029      (map (make-kalman-mapper state-cluster-prior)
1030          obn-clusters)))
1031
1032  #_(async-randomized-scan obn-clusters
1033      (make-kalman-mapper state-cluster-prior)
1034      clojure.pprint/pprint)
1035
1036  (let [accumulator (make-sow-reap)]
1037      (async-randomized-scan obn-clusters
1038          (make-kalman-mapper state-cluster-prior)
1039          (accumulator ::sow))
1040      (last (accumulator ::reap))) )

```

## 10 OZ FOR VISUALIZATION

```

1041
1042  From https://github.com/metasoarous/oz/blob/master/examples/clojupyter-example.
1043  ipynb

```

```

1044  (require '[clojupyter.misc.helper :as helper])
1045  (helper/add-dependencies '[metasoarous/oz "1.6.0-alpha2"])
1046  (require '[oz.notebook.clojupyter :as oz])

```

### 10.1 DEFN PLAY DATA

```

1047
1048  (do (defn play-data [& names]
1049      (for [n names
1050          i (range 20)]
1051          {:time i :item n :quantity (+ (Math/pow (* i (count n)) 0.8) (rand-int (count n)))})
1052
1053      (def stacked-bar
1054          {:data {:values (play-data "munchkin" "witch" "dog" "lion" "tiger" "bear")}
1055           :mark "bar"
1056           :encoding {:x {:field "time"}
1057                     :y {:aggregate "sum"
1058                         :field "quantity"
1059                         :type "quantitative"}}})

```

```
1060         :color {:field "item"}})})
1061   (oz/view! stacked-bar) )

1062 ;; Create spec, then visualize
1063 (def spec
1064   {:data {:url "https://gist.githubusercontent.com/metasoarous/4e6f781d353322a44b9cd3e4597
1065     :mark "point"
1066     :encoding {
1067       :x {:field "Horsepower", :type "quantitative"}
1068       :y {:field "Miles_per_Gallon", :type "quantitative"}
1069       :color {:field "Origin", :type "nominal"}}})
1070   (oz/view! spec)

1071 (oz/view!
1072   [:div
1073     [:h1 "A little hiccup example"]
1074     [:p "Try drinking a glass of water with your head upside down"]
1075     [:div {:style {:display "flex" :flex-direction "row"}}
1076       [:vega-lite spec]
1077       [:vega-lite stacked-bar]]])
```

## 11 GAUSSIAN PROCESSES

The Extended Kalman Filter above is a generalization of linear regression.

### 11.1 RECURRENT LINEAR REGRESSION

Emacs 26.2 of 2019-04-12, org version: 9.2.2