

Composable Statistics

Brian Beckman

May 26, 2019

Contents

1	COMPOSABLE STATISTICS	2
1.1	TODO TANGLE, NOWEB	2
1.2	TODO GRAPHICS	2
1.3	CLOJURE	2
1.4	TODO HOW TO USE THIS DOCUMENT	2
2	INTRODUCTION	2
2.1	TODO: GENERATE NEW RANDOM DATA	3
3	RUNNING COUNT	3
3.1	THREAD-SAFE	3
3.2	AVOIDING <i>REDUCE</i>	4
4	RUNNING MEAN	4
4.1	REMOVING OUTPUT COUPLING	5
4.2	NUMERICAL CHECK	6
5	CORE.ASYNC	6
5.1	SHALLOW TUTORIAL	6
5.2	DEEP TUTORIAL	6
6	ASYNC DATA STREAMS	13
6.1	ASYNC RUNNING MEAN	15
7	RUNNING STDDEV	17
7.1	BRUTE-FORCE (SCALAR VERSION)	17
7.2	DEF Z2S: SMALLER EXAMPLE	18
7.3	REALLY DUMB RECURRENCE	18
7.4	SCHOOL VARIANCE	18
7.5	DEFN MAKE SCHOOL STATS MAPPER	18
7.6	DEFN MAKE RECURRENT STATS MAPPER	19
7.7	DEFN MAKE WELFORD'S STATS MAPPER	20
8	WINDOWED STATISTICS	20
8.1	DEF Z3S: MORE SAMPLE DATA	21
8.2	DEFN MAKE SLIDING STATS MAPPER	22

34	9 KALMAN FILTER	23
35	9.1 BASIC LINEAR ALGEBRA	23
36	9.2 DEFN SYMMETRIC PART	24
37	9.3 DEFN ANTI-SYMMETRIC PART	24
38	9.4 DEFN KALMAN UPDATE: GENERAL EXTENDED KALMAN FILTER	27
39	9.5 DEFN MAKE-KALMAN-MAPPER	29
40	10 VISUALIZATION SANDBOX	31
41	11 VISUALIZATION	31
42	11.1 CLJ-REFACTOR	31
43	11.2 INCANTER	31
44	11.3 OZ	32
45	12 GAUSSIAN PROCESSES	33
46	12.1 RECURRENT LINEAR REGRESSION	33
47	13 SANDBOX	33

48 1 COMPOSABLE STATISTICS

49 1.1 TODO TANGLE, NOWEB

50 This will become a fully literate program via org-babel `tangle` and `noweb`. At present, it is an org file
51 converted directly from an old iPython (Jupyter) notebook. Jupyter proved not to scale well.

52 1.2 TODO GRAPHICS

53 We still have to work out how to embed graphics in this file.

54 1.3 CLOJURE

55 We prefer Clojure to Python for this exercise due to Clojure's concurrency primitives, especially atoms and
56 `core.async`. Python is growing and improving rapidly, so we may return to it someday.

57 1.3.1 TODO PROJECT.CLJ

58 At present, the critical file `project.clj` is external to this document. It will be one of the first to tangle.

59 The best sites for learning Clojure by example are clojuredocs.org and 4clojure.org. A recommended
60 book is Clojure for the Brave and True.

61 1.4 TODO HOW TO USE THIS DOCUMENT

62 Explain how to run Clojure code inside an org-mode buffer, how to tangle and weave, etc.

- 63 1. Install leiningen <https://leiningen.org/> (this is all you need for Clojure)

64 2 INTRODUCTION

65 We want to compute descriptive statistics in constant memory. We want exactly the same code to run over
66 sequences distributed in space as runs over sequences distributed in time. Sequences distributed in space
67 are vectors, lists, arrays, lazy or not. Sequences distributed over time are asynchronous streams. Descriptive
68 statistics range from `count`, `mean`, `max`, `min`, and `variance` to Kalman filters and Gaussian processes. We
69 decouple computation from data delivery by packaging computation in composable functions.

70 Some sample scalar data:

```

71 (def zs [-0.178654, 0.828305, 0.0592247, -0.0121089, -1.48014,
72         -0.315044, -0.324796, -0.676357, 0.16301, -0.858164])

```

73 2.1 TODO: GENERATE NEW RANDOM DATA

74 3 RUNNING COUNT

75 The traditional and obvious way with `reduce` and `reductions` ([https://clojuredocs.org/clojure.](https://clojuredocs.org/clojure.core/reduce)
76 `core/reduce`). *Reduce* takes three arguments: a binary function, an initial value, and a space-sequence of
77 inputs.

```

78 (reduce
79   (fn [count datum] (inc count)) ; binary function
80   0                             ; initial value
81   zs)                            ; space sequence

```

```

82 10

```

83 ... with all intermediate results:

```

84 (reductions (fn [c z] (inc c)) 0 zs)
85
86           0  1  2  3  4  5  6  7  8  9 10

```

86 3.1 THREAD-SAFE

87 Overkill for sequences in space, but safe for multiple threads from asynchronous streams. It also shows
88 (1) *let-over-lambda* (LOL): closing over mutable state variables, and (2) transactional mutation, i.e., *atomic*
89 *updates*. LOL is semantically equivalent to data encapsulation in OOP, and transactions are easier to verify
90 than is OOP with locks and mutexes.

91 The following has a defect: we need `initial-count` both to initialize the atom and to initialize the
92 `reduce` call. This defect must be traded off against the generalizable form or *functional type* of the reducible,
93 namely (estimate, measurement) → estimate. We get rid of this defect later.

```

94 (let [initial-count 0] ; Must use this twice below.
95   (reduce
96     ; Let-over-lambda (anonymous "object") follows.
97     ; "Atom" is a transactional (thread-safe) type in Clojure.
98     (let [running-count (atom initial-count)]
99       ; That was the "let" of "LOL." Here comes the lambda:
100      ; Reducible closure over "running-count."
101      (fn [c z] ; Here's the "lambda" of "LOL"
102        (swap! running-count inc) ; transactional update
103        @running-count))
104     ; safe "read" of the atom ~> new value for c
105     initial-count
106     zs))

```

```

107 10

```

108 Showing all intermediate results:

```

109 (let [initial-count 0]
110   (reductions ; <-- this is the only difference to above
111     (let [running-count (atom initial-count)]
112       (fn [c z]

```

```

113         (swap! running-count inc)
114         @running-count)) ; ~~> new value for c
115     initial-count
116     zs))
117
118     0 1 2 3 4 5 6 7 8 9 10

```

118 3.2 AVOIDING REDUCE

119 Reduce only works in space, not in time. Avoiding reduce decouples the statistics code (“business logic”) from the space environment (“plumbing”). That spaces environment delivers data from vectors, lists, etc.). We want to be able to switch out an environment that delivers data from space for an environment that delivers data points *z* from time.

123 The following is a thread-safe LOL, without `reduce`. We *map* the LOL over a space-sequence in memory to produce exactly the same result as with `reduce`. The mappable LOL does not need an accumulator argument for `count`.

126 Below, we map *exactly* the same mappable LOL over asynchronous streams.

127 A subtle defect: the output is still coupled to the computing environment through `print`. We get rid of that, too, below.

```

129 (dorun ; <-- Discard 'nil's produced by "print."
130     (map
131         (let [running-count (atom 0)]
132             (fn [z] ; <-- one fewer argument
133                 (swap! running-count inc)
134                 (print (str @running-count " "))))
135         zs))
136
137     1 2 3 4 5 6 7 8 9 10

```

137 4 RUNNING MEAN

138 Consider the following general scheme for recurrence: *a new statistic is an old statistic plus a correction*.

139 The *correction* is a *gain* times a *residual*. For running mean, the residual is the difference between the new measurement *z* and the old mean *x*. The gain is $1/(n+1)$, where *n* is *count-so-far*. *n* is a statistic, too, so it is an *old* value, computed and saved before the current observation *z* arrived.

142 /The correction therefore depends only on the new input *z* and on old statistics *x* and *n*. The correction does not depend on new statistics/.

144 Mathematically, write the general recurrence idea without subscripts as

$$x \leftarrow x + K(z - x)$$

145 or, with Lamport’s notation, wherein new versions of old values get a prime, as an equation

$$x' = x + K(z - x)$$

146 (*z* does not have a prime; it is the only exception to the rule that new versions of old quantities have primes).

148 Contrast the noisy traditional form, which introduces another variable, the index *n*. This traditional form is objectively more complicated than either of the two above:

$$x_{n+1} = x_n + K(n)(z_{n+1} - x_n)$$

```

150 (dorun
151     (map
152         (let [running-stats (atom {:count 0, :mean 0})]

```

```

153         (fn [z]
154           (let [{x :mean, n :count} @running-stats
155                 n+1 (inc n) ; cool variable name!
156                 K (/ 1.0 n+1)]
157             (swap! running-stats conj
158                   [:count n+1]
159                   [:mean (+ x (* K (- z x)))]))
160             (println @running-stats)))
161         zs))
162
163 { :count 1, :mean -0.178654}
164 { :count 2, :mean 0.3248255}
165 { :count 3, :mean 0.2362919}
166 { :count 4, :mean 0.1741917}
167 { :count 5, :mean -0.156674640000000003}
168 { :count 6, :mean -0.18306953333333337}
169 { :count 7, :mean -0.20331617142857145}
170 { :count 8, :mean -0.262446275}
171 { :count 9, :mean -0.21517335555555556}
172 { :count 10, :mean -0.27947242}

```

The swap above calls conj on the current contents of the atom running-stats and on the rest of the arguments, namely [:count n+1, :mean ...]. conj is the idiom for “updating” a hashmap, the hashmap in the atom, the hashmap that starts off as {:count 0, :mean 0}.

4.1 REMOVING OUTPUT COUPLING

Remove println from inside the LOL function of z. Now the LOL function of z is completely decoupled from its environment. Also, abstract a “factory” method for the LOL, *make-running-stats-mapper*, to clean up the line that does the printing.

4.1.1 MAKE-RUNNING-STATS-MAPPER

```

180 (defn make-running-stats-mapper []
181   (let [running-stats (atom {:count 0 :mean 0 :datum 0})]
182     (fn [z]
183       (let [{x :mean, n :count, _ :datum} @running-stats
184             n+1 (inc n)
185             K (/ 1.0 n+1)]
186         (swap! running-stats conj
187               [:count n+1]
188               [:mean (+ x (* K (- z x)))]
189               [:datum z]))
190         @running-stats)))
191
192 (clojure.pprint/pprint (map (make-running-stats-mapper) zs))
193
194 { :count 1, :mean -0.178654, :datum -0.178654}
195 { :count 2, :mean 0.3248255, :datum 0.828305}
196 { :count 3, :mean 0.2362919, :datum 0.0592247}
197 { :count 4, :mean 0.1741917, :datum -0.0121089}
198 { :count 5, :mean -0.156674640000000003, :datum -1.48014}
199 { :count 6, :mean -0.18306953333333337, :datum -0.315044}
200 { :count 7, :mean -0.20331617142857145, :datum -0.324796}
201 { :count 8, :mean -0.262446275, :datum -0.676357}
202 { :count 9, :mean -0.21517335555555556, :datum 0.16301}
203 { :count 10, :mean -0.27947242, :datum -0.858164}

```

4.2 NUMERICAL CHECK

The last value of the running mean is $-0.279...42$. Check that against an independent calculation.

1. DEFN MEAN

```
(defn mean [zs] (/ (reduce + zs) (count zs)))
(println (mean zs))

-0.27947242
```

5 CORE.ASYNC

For data distributed over time, we'll use Clojure's `core.async`. `Core.async` has some subtleties that we analyze below.

```
(require
 '[clojure.core.async
  :refer
  [sliding-buffer dropping-buffer buffer
   <!! <!, >!, >!!
   go chan onto-chan close!
   thread alts! alts!! timeout]])
```

5.1 SHALLOW TUTORIAL

<https://github.com/clojure/core.async/blob/master/examples/walkthrough.clj>

5.2 DEEP TUTORIAL

The asynchronous, singleton `go` thread is loaded with very lightweight *pseudothreads* (my terminology, not standard; most things you will read or see about Clojure.async does not carefully distinguish between threads and pseudothreads, and I think that's not helpful).

Pseudothreads are lightweight state machines that pick up where they left off. It is feasible to have thousands, even millions of them. Pseudothreads don't block, they *park*. *Parking* and *unparking* are very fast. We can write clean code with pseudothreads because our code looks like it's blocked waiting for input or blocked waiting for buffer space. Code with blocking I/O is easy to write and to understand. Code in `go` forms doesn't actually block, just looks like it.

Some details are tricky and definitely not easy to divine from the documentation. Hickey's video from InfoQ 2013 (<https://www.infoq.com/presentations/core-async-clojure>) is more helpful, but you can only appreciate the fine points after you've stumbled a bit. I stumbled over the fact that buffered and unbuffered channels have different synchronization semantics. Syntactically, they look the same, but you cannot, in general, run the same code over an unbuffered channel that works on a buffered channel. Hickey says this, but doesn't nail it to the mast; doesn't emphasize it with an example, as I do here in this deep tutorial. He motivates the entire library with the benefits of first-class queues, but fails to emphasize that, by default, a channel is not a queue but a blocking rendezvous. He does mention it, but one cannot fully appreciate the ramifications from a passing glance.

5.2.1 COMMUNICATING BETWEEN THREADS AND PSEUDOTHEADS

Write output to unbuffered channel `c` via `>!` on the asynchronous `go` real-thread and read input from the same channel `c` via `<!!` on the UI/REPL `println` real-thread. We'll see later that writing via `>!!` to an unbuffered channel blocks the UI real-thread, so we can't write before reading unbuffered on the UI/REPL real-thread. However, we can write before reading on a non-blocking pseudothread, and no buffer space is needed.

```

245 (let [c (chan)]           ;; unbuffered chan
246     (go (>! c 42))       ;; parks if no space in chan
247     (println (<!! c))    ;; blocks UI/REPL until data on c
248     (close! c))          ;; idiom; may be harmless overkill

```

```

249 42

```

250 In general, single-bang forms work on `go` pseudothreads, and double-bang forms work on real, heavy-weight, Java threads like the UI/REPL thread behind this notebook. In the rest of this notebook, “thread” means “real thread” and we write “pseudothread” explicitly when that’s what we mean.

253 I don’t address thread leakage carefully in this tutorial, mostly because I don’t yet understand it well. I may overkill by closing channels redundantly.

255 5.2.2 CHANNEL VOODOO FIRST

256 Writing before reading seems very reasonable, but it does not work on unbuffered channels, as we see below. Before going there, however, let’s understand more corners of the example above.

258 The `go` form itself returns a channel:

```

259 (clojure.repl/doc go)

```

```

260 -----
261 clojure.core.async/go
262 ([& body])
263 Macro
264   Asynchronously executes the body, returning immediately to the
265   calling thread. Additionally, any visible calls to <!, >! and alt!/alts!
266   channel operations within the body will block (if necessary) by
267   ‘parking’ the calling thread rather than tying up an OS thread (or
268   the only JS thread when in ClojureScript). Upon completion of the
269   operation, the body will be resumed.
270
271   Returns a channel which will receive the result of the body when
272   completed

```

273 I believe “the calling thread” above refers to a pseudothread inside the `go` real-thread, but I am not sure because of the ambiguities in the official documentation between “blocking” and “parking” and between “thread” and “well, we don’t have a name for them, but Brian calls them ‘pseudothreads’.”

276 Is the channel returned by `go` the same channel as `c`?

```

277 (let [c (chan)]
278     (println {:c-channel c})
279     (println {:go-channel (go (>! c 42))})
280     (println {:c-coughs-up (<!! c)})
281     (println {:close-c (close! c)}))

```

```

282 {:c-channel #object[clojure.core.async.impl.channels.ManyToManyChannel 0x3ee3f076 clojure.
283 {:go-channel #object[clojure.core.async.impl.channels.ManyToManyChannel 0x443b5535 clojure.
284 {:c-coughs-up 42}
285 {:close-c nil}

```

286 No, `c` is a different channel from the one returned by `go`. Consult the documentation for `go` once more:

```

287 (clojure.repl/doc go)

```

```

288 -----
289 clojure.core.async/go
290 ([& body])
291 Macro
292   Asynchronously executes the body, returning immediately to the
293   calling thread. Additionally, any visible calls to <!, >! and alt!/alts!
294   channel operations within the body will block (if necessary) by
295   'parking' the calling thread rather than tying up an OS thread (or
296   the only JS thread when in ClojureScript). Upon completion of the
297   operation, the body will be resumed.
298
299   Returns a channel which will receive the result of the body when
300   completed

```

301 We should be able to read from the channel returned by go; call it d:

```

302 (let [c (chan)
303       d (go (>! c 42))] ;; 'let' in Clojure is sequential,
304                       ;; like 'let*' in Scheme or Common Lisp,
305                       ;; so 'd' has a value, here.
306   (println {:c-coughs-up (<!! c),   ;; won't block
307             :d-coughs-up (<!! d)}) ;; won't block
308   (close! c)
309   (close! d))
310 {:c-coughs-up 42, :d-coughs-up true}

```

311 d's coughing up true means that the body of the go, namely (>! c 42) must have returned true,
 312 because d coughs up "the result of the body when completed." Let's see whether our deduction matches
 313 documentation for >!:

```

314 (clojure.repl/doc >!)
315 -----
316 clojure.core.async/>!
317 ([port val])
318   puts a val into port. nil values are not allowed. Must be called
319   inside a (go ...) block. Will park if no buffer space is available.
320   Returns true unless port is already closed.

```

321 Sure enough. But something important is true and not obvious from this documentation. Writing to c
 322 inside the go block parks the pseudothread because no buffer space is available: c was created with a call to
 323 chan with no arguments, so no buffer space is allocated. Only when reading from c does the pseudothread
 324 unpark. How? There is no buffer space. Reading on the UI thread manages to short-circuit any need for a
 325 buffer and unpark the pseudothread. Such short-circuiting is called a *rendezvous* in the ancient literature of
 326 concurrency. Would the pseudothread unpark if we read inside a go block and not on the UI thread?

```

327 (let [c (chan)
328       d (go (>! c 42))
329       e (go (<! c))]
330   (clojure.pprint/pprint {
331     :c-channel c, :d-channel d, :e-channel e,
332     :e-coughs-up (<!! e),   ;; won't block
333     :d-coughs-up (<!! d)}) ;; won't block
334   (close! c)
335   (close! d)
336   (close! e))

```



```

337 {:c-channel
338  #object[clojure.core.async.impl.channels.ManyToManyChannel 0x3c39496e "clojure.core.async."
339  :d-channel
340  #object[clojure.core.async.impl.channels.ManyToManyChannel 0x2857d07d "clojure.core.async."
341  :e-channel
342  #object[clojure.core.async.impl.channels.ManyToManyChannel 0xa0a1d50 "clojure.core.async."
343  :e-coughs-up 42,
344  :d-coughs-up true}

```

345 Yes, the pseudothread that parked when 42 is put on `c` via `>!` unparks when 42 is taken off via `<!`.
 346 Channel `d` represents the parking step and channel `e` represents the unparking step. All three channels are
 347 different.

348 So now we know how to short-circuit or rendezvous unbuffered channels. In fact, the order of reading
 349 and writing (taking and putting) does not matter in the nebulous, asynchronous world of pseudothreads.
 350 How Einsteinian is that? The following takes (reads) from `c` on `e` before putting (writing) to `c` on `d`. That's
 351 the same as above, only in the opposite order.

```

352 (let [c (chan)
353       e (go (<! c))
354       d (go (>! c 42))]
355   (clojure.pprint/pprint {
356     :c-channel c, :d-channel d, :e-channel e,
357     :e-coughs-up (<!! e), ;; won't block
358     :d-coughs-up (<!! d)}) ;; won't block
359   (close! c)
360   (close! d)
361   (close! e))

362 {:c-channel
363  #object[clojure.core.async.impl.channels.ManyToManyChannel 0x6f9726b1 "clojure.core.async."
364  :d-channel
365  #object[clojure.core.async.impl.channels.ManyToManyChannel 0x7448570d "clojure.core.async."
366  :e-channel
367  #object[clojure.core.async.impl.channels.ManyToManyChannel 0x6cad36bb "clojure.core.async."
368  :e-coughs-up 42,
369  :d-coughs-up true}

```

370 5.2.3 PUTS BEFORE TAKES CONSIDERED RISKY

371 `>!!`, by default, blocks if called too early on an unbuffered real thread. We saw above that parked pseu-
 372 dothreads don't block: you can read and write to channels in `go` blocks in any order. However, that's not
 373 true with threads that actually block. The documentation is obscure, though not incorrect, about this fact.

```

374 (clojure.repl/doc >!!)

375 -----
376 clojure.core.async/>!!
377 ([port val])
378   puts a val into port. nil values are not allowed. Will block if no
379   buffer space is available. Returns true unless port is already closed.

```

380 When is “no buffer space available?” It turns out that the default channel constructor makes a channel
 381 with no buffer space allocated by default.

```

382 (clojure.repl/doc chan)

```

```

-----
clojure.core.async/chan
([] [buf-or-n] [buf-or-n xform] [buf-or-n xform ex-handler])
  Creates a channel with an optional buffer, an optional transducer
  (like (map f), (filter p) etc or a composition thereof), and an
  optional exception-handler. If buf-or-n is a number, will create
  and use a fixed buffer of that size. If a transducer is supplied a
  buffer must be specified. ex-handler must be a fn of one argument -
  if an exception occurs during transformation it will be called with
  the Throwable as an argument, and any non-nil return value will be
  placed in the channel.

```

We can test the blocking-on-unbuffered case as follows. The following code will block at the line (`>!! c 42`), as you'll find if you uncomment the code (remove `#_` at the beginning) and run it. You'll have to interrupt the Kernel using the "Kernel" menu at the top of the notebook, and you might have to restart the Kernel, but you should try it once.

```

#_(let [c (chan)]
      (>!! c 42)
      (println (<!! c))
      (close! c))

```

The following variation works fine because we made "buffer space" before writing to the channel. The only difference to the above is the 1 argument to the call of `chan`.

```

(let [c (chan 1)]
  (>!! c 42)
  (println (<!! c))
  (close! c))

```

42

The difference between the semantics of the prior two examples is not subtle: one hangs the kernel and the other does not. However, the difference in the syntax is subtle and easy to miss.

We can read on the asynchronous `go` pool from the buffered channel `c` because the buffered write (`>!! c`) on the UI thread doesn't block:

```

(let [c (chan 1)]
  (>!! c 42)
  (println {:go-channel-coughs-up (<!! (go (<! c)))})
  (close! c))

```

```

{:go-channel-coughs-up 42}

```

1. ORDER DOESN'T MATTER, SOMETIMES

We can do things backwards, reading before writing, even without a buffer. Read from channel (`<! c`) on the async `go` thread "before" writing to (`>!! c 42`) on the REPL / UI thread. "Before," here, of course, means syntactically or lexically "before," not temporally.

```

(let [c (chan) ;; NO BUFFER!
      d (go (<! c)) ;; park a pseudothread to read c
      e (>!! c 42)] ;; blocking write unparks c's pseudothread
  (println {:c-hangs '(<!! c),
            :d-coughs-up (<!! d),
            :what's-e e})
  (close! c) (close! d))

```

```
{:c-hangs (<!! c), :d-coughs-up 42, :what's-e true}
```

Why did `>!!` produce `true`? Look at docs again:

```
(clojure.repl/doc >!!)
```

```
-----
clojure.core.async/>!!
([port val])
  puts a val into port. nil values are not allowed. Will block if no
  buffer space is available. Returns true unless port is already closed.
```

Ok, now I fault the documentation. `>!!` will block if there is no buffer space available *and* if there is no *rendezvous* available, that is, no pseudothread parked waiting for `<!!`. I have an open question in the Google group for Clojure about this issue with the documentation.

To get the value written in into `c`, we must read `d`. If we tried to read it from `c`, we would block forever because `>!!` blocks when there is no buffer space, and `c` never has buffer space. We get the value out of the `go` nebula by short-circuiting the buffer, by a rendezvous, as explained above.

`e's` being `true` means that `c` wasn't closed. `(>!! c 42)` should hang.

```
(let [c (chan) ;; NO BUFFER!
      d (go (<!! c)) ;; park a pseudothread to read c
      e (>!! c 42) ;; blocking write unparks c's pseudothread
      f '(hangs (>!! c 43))] ;; is 'c' closed?
  (println {:c-coughs-up '(hangs (<!! c)),
            :d-coughs-up (<!! d),
            :what's-e e,
            :what's-f f})
  (close! c) (close! d))
```

```
{:c-coughs-up (hangs (<!! c)), :d-coughs-up 42, :what's-e true, :what's-f (hangs (>!! c 43))}
```

StackOverflow reveals a way to find out whether a channel is closed by peeking under the covers (<https://stackoverflow.com/questions/24912971>):

```
(let [c (chan) ;; NO BUFFER!
      d (go (<!! c)) ;; park a pseudothread to read c
      e (>!! c 42) ;; blocking write unparks c's pseudothread
      f (clojure.core.async.impl.protocols/closed? c)]
  (println {:c-coughs-up '(hangs (<!! c)),
            :d-coughs-up (<!! d),
            :c-is-open-at-e? e,
            :c-is-open-at-f? f})
  (close! c) (close! d))
```

```
{:c-coughs-up (hangs (<!! c)), :d-coughs-up 42, :c-is-open-at-e? true, :c-is-open-at-f? false}
```

2. ORDER DOES MATTER, SOMETIMES

Order does matter this time: Writing blocks the UI thread without a buffer and no parked read (rendezvous) in the `go` nebula beforehand. I hope you can predict that the following will block even before you run it. To be sure, run it, but you'll have to interrupt the kernel as before.

```

470   #_(let [c (chan)
471         e (>!! c 42) ;; blocks forever
472         d (go (<! c))])
473   (println {:c-coughs-up '(this will hang (<!! c)),
474            :d-coughs-up (<!! d),
475            :what's-e     e})
476   (close! c) (close! d))

```

477 5.2.4 TIMEOUTS: DON'T BLOCK FOREVER

478 In all cases, blocking calls like `>!!` to unbuffered channels without timeout must appear *last* on the UI,
 479 non-`go`, thread, and then only if there is some parked pseudothread that's waiting to read the channel by
 480 short-circuit (rendezvous). If we block too early, we won't get to the line that launches the async `go` nebula
 481 and parks the short-circuitable pseudothread—parks the rendezvous.

482 The UI thread won't block forever if we add a timeout. `alts!!` is a way to do that. The documen-
 483 tation and examples are difficult, but, loosely quoting (emphasis and edits are mine, major ones in square
 484 brackets):

```

485   (alts!! ports & {:as opts})

```

486 This destructures all keyword options into `opts`. We don't need `opts` or the `:as` keyword below.

487 Completes at most one of several channel operations. [/Not for use inside a (go ...) block./]
 488 **ports is a vector of channel endpoints**, [A channel endpoint is] either a channel to take from
 489 or a vector of [channel-to-put-to val-to-put] pairs, in any combination. Takes will be
 490 made as if by `<!!`, and puts will be made as if by `>!!`. If more than one port operation is ready,
 491 a non-deterministic choice will be made unless the `:priority` option is true. If no operation
 492 is ready and a `:default` value is supplied, [=default-val :default=] will be returned, otherwise
 493 `alts!!` will [/block/ xxxpark ?] until the first operation to become ready completes. **Re-**
 494 **turns [val port] of the completed operation**, where `val` is the value taken for takes, and a
 495 boolean (`true` unless already closed, as per `put!`) for puts. `opts` are passed as `:key val...`
 496 Supported options: `:default val` - the value to use if none of the operations are immediately
 497 ready `:priority true` - (default `nil`) when `true`, the operations will be tried in order. Note:
 498 there is no guarantee that the port exprs or val exprs will be used, nor in what order should they
 499 be, so they should not be depended upon for side effects.

```

500   (alts!! ...) returns a [val port] 2-vector.
501   (second (alts!! ...)) is a wrapper of channel c We can't write to the resulting timeout chan-
502   nel because we didn't give it a name.

```

503 That's a lot of stuff, but we can divine an idiom: pair a channel `c` that *might* block with a fresh `timeout`
 504 channel in an `alts!!`. At most one will complete. If `c` blocks, the `timeout` will cough up. If `c` coughs up
 505 before the `timeout` expires, the `timeout` quietly dies (question, is it closed? Will it be left open and leak?)

506 For a first example, let's make a buffered thread that won't block and pair it with a long timeout. You
 507 will see that it's OK to write 43 into this channel (the `[c 43]` term is an implied write; that's clear from
 508 the documentation). `c` won't block because it's buffered, it returns immediately, long before the `timeout`
 509 could expire.

```

510   (let [c (chan 1)
511         a (alts!! ; outputs a [val port] pair; throw away the val
512               ; here are the two channels for `alts!!`
513               [[c 43] (timeout 2500)])])
514   (clojure.pprint/pprint {:c c, :a a})
515   (let [d (go (<! c))])
516   (println {:d-returns (<!! d)})
517   (close! c)

```

```

518 {:c
519 #object[clojure.core.async.impl.channels.ManyToManyChannel 0x44e571a7 "clojure.core.async
520 :a
521 [true
522 #object[clojure.core.async.impl.channels.ManyToManyChannel 0x44e571a7 "clojure.core.async
523 {:d-returns 43}]

```

But, if we take away the buffer, the timeout channel wins. The only difference to the above is that instead of creating `c` via `(chan 1)`, that is, with a buffer of length 1, we create it with no buffer (and we quoted out the blocking read of `d` with a tick mark).

```

527 (let [c (chan)
528       a (alts!! ; outputs a [val port] pair; throw away the val
529             ; here are the two channels for `alts!!`
530           [[c 43] (timeout 2500)])]
531   (clojure.pprint/pprint {:c c, :a a})
532   (let [d (go (<! c))]
533     (println {:d-is d})
534     ' (println {:d-returns (<!! d)})) ; blocks
535   (close! c))

536 {:c
537 #object[clojure.core.async.impl.channels.ManyToManyChannel 0x7c666fe7 "clojure.core.async
538 :a
539 [nil
540 #object[clojure.core.async.impl.channels.ManyToManyChannel 0x7a5f1107 "clojure.core.async
541 {:d-is #object[clojure.core.async.impl.channels.ManyToManyChannel 0x65579a85 clojure.core.

```

6 ASYNC DATA STREAMS

The following writes at random times (`>!`) to a parking channel `echo-chan` on an `async go` fast pseudothread. The UI thread block-reads (`<!!`) some data from `echo-chan`. The UI thread leaves values in the channel and thus leaks the channel according to the documentation for `close!` here <https://clojure.github.io/core.async/api-index.html#C>. To prevent the leak permanently, we close the channel explicitly.

```

548 (def echo-chan (chan))
549
550 (doseq [z zs] (go (Thread/sleep (rand 100)) (>! echo-chan z)))
551 (dotimes [_ 3] (println (<!! echo-chan)))
552
553 (println {:echo-chan-closed?
554           (clojure.core.async.impl.protocols/closed? echo-chan)})
555 (close! echo-chan)
556 (println {:echo-chan-closed?
557           (clojure.core.async.impl.protocols/closed? echo-chan)})

558 -0.676357
559 -0.324796
560 -0.0121089
561 {:echo-chan-closed? false}
562 {:echo-chan-closed? true}

```

We can chain channels, again with leaks that we explicitly close. Also, we must not `>!` (send) a nil to `repl-chan`, and `<!` can produce nil from `echo-chan` after the timeout and we close `echo-chan`.

```

565 (clojure.repl/doc <!)
566 -----
567 clojure.core.async/<!
568 ([port])
569   takes a val from port. Must be called inside a (go ...) block. Will
570   return nil if closed. Will park if nothing is available.
571
572   Every time you run the block of code below, you will probably get a different result, by design.
573
574 (def echo-chan (chan))
575 (def repl-chan (chan))
576
577 ;; >! chokes on nulls. <! echo-chan can cough up nil if we time out
578 ;; and close the channel. The following line will throw an exception
579 ;; unless we don't close the channel at the end of this code-block.
580
581 ;; (dotimes [_ 10] (go (>! repl-chan (<! echo-chan))))
582
583 ;; Instead of throwing an exception, just put a random character
584 ;; like \? down the pipe after the echo-chan is closed:
585
586 (dotimes [_ 10] (go (>! repl-chan (or (<! echo-chan) \?))))
587
588 (doseq [z zs] (go (Thread/sleep (rand 100)) (>! echo-chan z)))
589
590 (dotimes [_ 3]
591   (println (<!! (second (alts!! [repl-chan
592                                   (timeout 500)])))))
593
594 ;; Alternatively, we can avoid the exception by NOT closing echo-chan.
595 ;; Not closing echo chan will leak it, and that's a lousy idea.
596
597 (close! echo-chan)
598
599 (close! repl-chan)
600
601 -1.48014
602 -0.178654
603 0.16301

```

601 Reading from echo-chan may hang the UI thread because the UI thread races the internal go thread
602 that reads echo-chan, but the timeout trick works here as above.

```

603 (def echo-chan (chan))
604 (def repl-chan (chan))
605
606 (dotimes [_ 10] (go (>! repl-chan (or (<! echo-chan) \?))))
607 (doseq [z zs] (go (Thread/sleep (rand 100)) (>! echo-chan z)))
608 (dotimes [_ 3]
609   (println (<!! (second (alts!! [echo-chan
610                                   (timeout 500)])))))
611
612 (close! echo-chan)
613 (close! repl-chan)

```

```

614 nil
615 nil
616 nil

```

617 `println` on a `go` pseudoprocess works if we wait long enough. This, of course, is bad practice or “code
618 smell.”

```

619 (def echo-chan (chan))
620
621 (doseq [z zs] (go (Thread/sleep (rand 100)) (>! echo-chan z)))
622 (dotimes [_ 3] (go (println (<! echo-chan)))))
623
624 (Thread/sleep 500) ; no visible output if you remove this line.
625 (close! echo-chan)

626 -0.178654
627 0.828305
628 -0.315044

```

629 6.1 ASYNC RUNNING MEAN

630 6.1.1 DEFN ASYNC-RANDOMIZED-SCAN

631 We want `running-stats` called at random times and with data in random order. A *transducer*, (`map`
632 `mapper`), lets us collect items off the buffer. The size of the buffer does not matter, but we must specify
633 it. Notice that the side-effector `effector` is passed in, so `async-randomized-scan` remains decoupled
634 from its environment.

635 In this style of programming, the asynchronous stream might sometimes be called a *functor*, which is
636 anything that’s mappable, anything you can map over.

```

637 (defn async-randomized-scan [zs mapper effector]
638   (let [transducer (map mapper)
639         ; give buffer length if there is a transducer
640         echo-chan (chan (buffer 1) transducer)]
641     (doseq [z zs]
642       (go (Thread/sleep (rand 100)) (>! echo-chan z)))
643     (dotimes [_ (count zs)] (effector (<!! echo-chan)))
644     (close! echo-chan)))
645
646 (async-randomized-scan zs (make-running-stats-mapper) println)

647 {:count 1, :mean -0.324796, :datum -0.324796}
648 {:count 2, :mean -0.9024679999999999, :datum -1.48014}
649 {:count 3, :mean -0.6611966666666667, :datum -0.178654}
650 {:count 4, :mean -0.5746585, :datum -0.315044}
651 {:count 5, :mean -0.6313596, :datum -0.858164}
652 {:count 6, :mean -0.5162622166666667, :datum 0.0592247}
653 {:count 7, :mean -0.41922332857142863, :datum 0.16301}
654 {:count 8, :mean -0.26328228750000005, :datum 0.828305}
655 {:count 9, :mean -0.23537413333333337, :datum -0.0121089}
656 {:count 10, :mean -0.27947242000000005, :datum -0.676357}

```

657 We don’t need to explicitly say `buffer`, but I prefer to do.

6.1.2 DEFN MAKE SOW REAP

The `effector` above just prints to the console. Suppose we want to save the data?

The following is a version of Wolfram's `Sow` and `Reap` that does not include tags. It uses `atom` for an effectful store because a `let` variable like `result` is not a `var` and `alter-var-root` won't work on `(let [result []] ...)`. An `atom` might be overkill.

`make-sow-reap` returns a message dispatcher in the style of *The Little Schemer*. It responds to namespaced keywords `::sow` and `::reap`. In the case of `::sow`, it returns an `effector` function that `conj`'s its input to the internal result atomically. In the case of `::reap`, it returns the value of the result accumulated so-far.

```
(do (defn make-sow-reap []
      (let [result (atom [])]
        (fn [msg]
          (cond
            (identical? msg ::sow)
            (fn [x] (swap! result #(conj % x)))
            (identical? msg ::reap)
            @result))))

      (let [accumulator (make-sow-reap)]
        (async-randomized-scan zs
                               (make-running-stats-mapper)
                               (accumulator ::sow))
        (last (accumulator ::reap)))

      :count 10 :mean -0.27947242 :datum -0.858164
```

Occasionally, there is some floating-point noise in the very low digits of the mean because `async-randomized-scan` scrambles the order of the inputs. The mean should always be almost equal to `-0.27947242`.

6.1.3 DEFN ASYNC NON RANDOM SCAN

Of course, the mean of any permutation of the data `zs` is the same, so the order in which data arrive does not change the final result, except for some occasional floating-point noise as mentioned above.

```
(do (defn async-non-random-scan [zs mapper effector]
      (let [transducer (map mapper)
            echo-chan (chan (buffer 1) transducer)]
        (go (doseq [z zs] (>! echo-chan z)))
        (dotimes [_ (count zs)] (effector (<!! echo-chan)))
        (close! echo-chan)))

      (let [accumulator (make-sow-reap)]
        (async-non-random-scan zs (make-running-stats-mapper)
                               (accumulator ::sow))
        (last (accumulator ::reap)))

      :count 10 :mean -0.27947242 :datum -0.858164
```

6.1.4 DEFN SYNC SCAN: WITH TRANSDUCER

Here is the modern way, with `transduce`, to reduce over a sequence of data, in order. It's equivalent to the non-random async version above. The documentation for `transduce` writes its parameters as `xform f coll`, and then says

reduce with a transformation of `f (xf)`. If `init` is not supplied, `(f)` will be called to produce it.

Our `xform` is `transducer`, or `(map mapper)`, and our `f` is `conj`, so this is an idiom for mapping because `(conj)`, with no arguments, returns `[]`, an appropriate `init`.

```
(do (defn sync-scan [zs mapper]
      (let [transducer (map mapper)]
        (transduce transducer conj zs)))
    (last (sync-scan zs (make-running-stats-mapper)))
    :count 10 :mean -0.27947242 :datum -0.858164)
```

We now have complete symmetry between space and time, space represented by the vector `zs` and time represented by values on `echo-chan` in random and in non-random order.

7 RUNNING STDDEV

7.1 BRUTE-FORCE (SCALAR VERSION)

The definition of variance is the following, for $N > 1$:

$$\frac{1}{N-1} \sum_{i=1}^N (z_i - \bar{z}_N)^2$$

The sum is the *sum of squared residuals*. Each residual is the difference between the i -th datum z_i and the mean \bar{z}_N of all N data in the sample. The outer constant, $1/(N-1)$ is Bessel's correction.

7.1.1 DEFN SSR: SUM OF SQUARED RESIDUALS

The following is *brute-force* in the sense that it requires all data up-front so that it can calculate the mean.

```
(do (defn ssr [sequ]
      (let [m (mean sequ)]
        (reduce #(+ %1 (* (- %2 m) (- %2 m)))
                0 sequ)))
    (ssr zs)
    3.5566483654807355)
```

7.1.2 DEFN VARIANCE

Call `ssr` to compute variance:

```
(do (defn variance [sequ]
      (let [n (count sequ)]
        (case n
          0 0
          1 (first sequ)
          #_default (/ (ssr sequ) (- n 1.0)))))
    (variance zs)
    0.3951831517200817)
```

7.2 DEF Z2S: SMALLER EXAMPLE

Let's do a smaller example:

```
(do (def z2s [55. 89. 144.])
    (variance z2s) )
```

2017.0

7.3 REALLY DUMB RECURRENCE

Remember our general form for recurrences, $x \leftarrow x + K \times (z - x)$?

We can squeeze running variance into this form in a really dumb way. The following is really dumb because:

1. it requires the whole sequence up front, so it doesn't run in constant memory
2. the intermediate values are meaningless because they refer to the final mean and count, not to the intermediate ones

But, the final value is correct.

```
(do (reductions
    (let [m (mean z2s) ; uh-oh, we refer to _all_ the data ??
          c (count z2s)]
      (fn [var z] (+ var (let [r (- z m)] ; residual
                             (/ (* r r) (- c 1.0))))))
    0 z2s) )
```

That was so dumb that we won't bother with a thread-safe, stateful, or asynchronous form.

7.4 SCHOOL VARIANCE

For an easy, school-level exercise, prove the following equation:

$$\frac{1}{N-1} \sum_{i=1}^N (z_i - \bar{z}_N)^2 = \frac{1}{N-1} \left(\sum_{i=1}^N (z_i^2) - N \bar{z}_N^2 \right)$$

Instead of the sum of squared residuals, *ssr*, accumulate the sum of squares, *ssq*.

School variance is exposed to *catastrophic cancellation* because *ssq* grows quickly. We fix that defect below.

We see that something is not best with this form because we don't use the old variance to compute the new variance. We do better below.

Of course, the same mapper works synchronously and asynchronously.

7.5 DEFN MAKE SCHOOL STATS MAPPER

and test it both synchronously and asynchronously, randomized and not:

```
(defn make-school-stats-mapper []
  (let [running-stats (atom {:count 0, :mean 0,
                             :variance 0, :ssq 0})]
    (fn [z]
      (let [{x :mean, n :count, s :ssq} @running-stats
            n+1 (inc n)
            K (/ 1.0 n+1)
            r (- z x)
            x' (+ x (* K r))] ;; Isn't prime notation nice?
```

```

777         s' (+ s (* z z))]
778     (swap! running-stats conj
779         [:count      n+1]
780         [:mean       x' ]
781         [:ssq        s' ]
782         [:variance (/ (- s' (* n+1 x' x')) (max 1 n))]))
783     @running-stats)))
784
785 (clojure.pprint/pprint (sync-scan z2s (make-school-stats-mapper)))
786
787 (async-randomized-scan z2s (make-school-stats-mapper) println)
788
789 (async-non-random-scan z2s (make-school-stats-mapper) println)

```

```

790 [{:count 1, :mean 55.0, :variance 0.0, :ssq 3025.0}
791  {:count 2, :mean 72.0, :variance 578.0, :ssq 10946.0}
792  {:count 3, :mean 96.0, :variance 2017.0, :ssq 31682.0}]
793 [{:count 1, :mean 144.0, :variance 0.0, :ssq 20736.0}
794  {:count 2, :mean 116.5, :variance 1512.5, :ssq 28657.0}
795  {:count 3, :mean 96.0, :variance 2017.0, :ssq 31682.0}
796  {:count 1, :mean 55.0, :variance 0.0, :ssq 3025.0}
797  {:count 2, :mean 72.0, :variance 578.0, :ssq 10946.0}
798  {:count 3, :mean 96.0, :variance 2017.0, :ssq 31682.0}]

```

7.6 DEFN MAKE RECURRENT STATS MAPPER

We already know the recurrence for the mean:

$$x \leftarrow x + K \cdot (z - x) = x + \frac{1}{n+1}(z - x)$$

We want a recurrence with a similar form for the variance. It takes a little work to prove, but it's still a school-level exercise. K remains $1/(n+1)$, the value needed for the new mean. We could define a pair of gains, one for the mean and one for the variance, but it would be less pretty.

$$v \leftarrow \frac{(n-1)v + K n (z-x)^2}{\max(1, n)}$$

```

804 (defn make-recurrent-stats-mapper []
805   (let [running-stats (atom {:count 0, :mean 0,
806                             :variance 0})]
807     (fn [z]
808       (let [{x :mean, n :count, v :variance} @running-stats
809             n+1 (inc n)
810             K (/ 1.0 (inc n))
811             r (- z x)
812             x' (+ x (* K r))
813             ssr (+ (* (- n 1) v) ; old ssr is (* (- n 1) v)
814                   (* K n r r))]
815         (swap! running-stats conj
816             [:count      n+1]
817             [:mean       x' ]
818             [:variance (/ ssr (max 1 n))]))
819       @running-stats)))
820
821 (async-non-random-scan z2s (make-recurrent-stats-mapper) println)

```

```

822 {:count 1, :mean 55.0, :variance 0.0}
823 {:count 2, :mean 72.0, :variance 578.0}
824 {:count 3, :mean 96.0, :variance 2017.0}

```

825 7.7 DEFN MAKE WELFORD'S STATS MAPPER

826 The above is equivalent, algebraically and numerically, to Welford's famous recurrence for the sum of
827 squared residuals S . In recurrences, we want everything on the right-hand sides of equations or left arrows
828 to be old, *prior* statistics, except for the new observation / measurement / input z . Welford's requires
829 the new, *posterior* mean on the right-hand side, so it's not as elegant as our recurrence above. However, it is
830 easier to remember!

$$S \leftarrow S + (z - x_N)(z - x_{N+1}) = S + (z - x)(z - (x + K(z - x)))$$

```

831 (do (defn make-welfords-stats-mapper []
832       (let [running-stats (atom {:count 0, :mean 0, :variance 0})]
833         (fn [z]
834           (let [{x :mean, n :count, v :variance} @running-stats
835                 n+1 (inc n)
836                 K (/ 1.0 n+1)
837                 r (- z x)
838                 x' (+ x (* K r))
839                 ssr (+ (* (- n 1) v)
840                       ;; only difference to recurrent variance:
841                       (* (- z x) (- z x')))]
842             (swap! running-stats conj
843                   [:count n+1]
844                   [:mean x']
845                   [:variance (/ ssr (max 1 n))]))
846           @running-stats)))
847
848 (async-non-random-scan
849   z2s (make-welfords-stats-mapper) println)
850
850 {:count 1, :mean 55.0, :variance 0.0}
851 {:count 2, :mean 72.0, :variance 578.0}
852 {:count 3, :mean 96.0, :variance 2017.0}

```

853 8 WINDOWED STATISTICS

854 Suppose we want running statistics over a history of fixed, finite length. For example, suppose we have
855 $N = 10$ data and we want the statistics in a window of length $w = 3$ behind the current value, inclusively.
856 When the first datum arrives, the window and the total include one datum. The window overhangs the
857 left until the third datum. When the fourth datum arrives, the window contains three data and the total
858 contains four data. After the tenth datum, we may consider three more steps marching the window "off
859 the cliff" to the right. The following figure illustrates (the first row corresponds to $n = 0$, not to $n = 1$):

860 We won't derive the following formulas, but rather say that they have been vetted at least twice inde-
861 pendently (in a C program and in a Mathematica program). The following table shows a unit test that we
862 reproduce. The notation is explained after the table.

863 Denote prior statistics by plain variables like m and corresponding posteriors by the same variables
864 with primes like m' . The posteriors j and u do not have a prime.

variable	description
n	prior count of data points; equals 0 when considering the first point
z	current data point
w	fixed, constant, maximum width of window; $w \geq 1$
j	posterior number of points left of the window; $j \geq 0$
u	posterior number of points including z in the running window; $1 \leq u \leq w$
m	prior mean of all points, not including z
m'	posterior mean of all points including z
m_j	prior mean of points left of the window, lagging w behind m
m'_j	posterior mean of points left of the window
m'_w	posterior mean of points in the window, including the current point z
v	prior variance, not including z
v'	posterior variance of all points including z
v_j	prior variance of points left of the window, lagging w behind u_n
v'_j	posterior variance of points left of the window
v'_w	posterior variance of points within the window

The recurrences for m , v , m_j , and v_j have only priors (no primes) on their right-hand sides. The values of m_w and v_w are not recurrences because the non-primed versions do not appear on the right-hand sides of equations 10 and 13. Those equations are simply transformations of the posteriors (values with primes) m' , m'_j , v' , and v'_j .

$$\begin{aligned}
 j &= \max(0, n + 1 - w) \\
 u &= n - j + 1 \\
 m' &= m + \frac{z - m}{n + 1} \\
 m'_j &= \begin{cases} m_j + \frac{z_j - m_j}{j} & j > 0 \\ 0 & \text{otherwise} \end{cases} \\
 m'_w &= \frac{(n + 1) m' - j m'_j}{u} \\
 v' &= \frac{(n - 1) v + \frac{n}{n + 1} (z - m)^2}{\max(1, n)} \\
 v'_j &= \begin{cases} \frac{j - 2}{j - 1} v_j + \frac{1}{j} (z_j - m_j)^2 & j > 1 \\ 0 & \text{otherwise} \end{cases} \\
 v'_w &= \frac{n v' + (n - w) v'_j + (n + 1) m'^2 - j m_j'^2 - u m_w'^2}{\max(1, u - 1)}
 \end{aligned}$$

Here is sample data we can compare with the unit test above.

8.1 DEF Z3S: MORE SAMPLE DATA

```
(def z3s [0.857454, 0.312454, 0.705325, 0.8393630, 1.637810,
          0.699257, -0.340016, -0.213596, -0.0418609, 0.054705])
```

The best algorithm we have found for tracking historical data is to keep a FIFO queue in a Clojure *vector* of length w . This is still constant memory because it depends only on the length w of the window, not on the length of the data stream.

8.1.1 DEFN PUSH TO BACK

```
(defn push-to-back [item vek]
  (conj (vec (drop 1 vek)) item))
```

8.2 DEFN MAKE SLIDING STATS MAPPER

```

880
881 (defn make-sliding-stats-mapper [w]
882   (let [running-stats (atom {:n 0, :m 0, :v 0,
883                               :win (vec (repeat w 0)),
884                               :mw 0, :vw 0,
885                               :mj 0, :vj 0})]
886     (fn [z]
887       (let [{:keys [m n v win mj vj]} @running-stats
888             zj (first win)
889             win' (push-to-back z win)
890             n+1 (double (inc n))
891             n-1 (double (dec n))
892             K (/ 1.0 n+1)
893             Kv (* n K)
894             r (- z m)
895             j (max 0, (- n+1 w))
896             u (- n+1 j)
897             m' (+ m (* K r))
898             rj (- zj mj)
899             mj' (if (> j 0), (+ mj (/ rj j)), 0)
900             mw' (/ (- (* n+1 m') (* j mj')) u)
901             v' (/ (+ (* n-1 v) (* Kv r r))
902                  (max 1 n))
903             vj' (if (> j 1)
904                  (let [j21 (/ (- j 2.0)
905                                (- j 1.0))]
906                    (+ (* j21 vj)
907                      (/ (* rj rj) j)))
908                  0)
909             vw' (let [t1 (- (* n v')
910                             (* (- n w) vj'))
911                      t2 (- (* n+1 m' m')
912                             (* j mj' mj'))
913                      t3 (- (* u mw' mw'))]
914                   (/ (+ t1 t2 t3)
915                      (max 1 (- u 1))))
916             ]
917       (swap! running-stats conj
918         [:n n+1 ]
919         [:m m' ]
920         [:v v' ]
921         [:mj mj' ]
922         [:vj vj' ]
923         [:mw mw' ]
924         [:vw vw' ]
925         [:win win'])
926       @running-stats)))
927
928 (clojure.pprint/print-table
929   [:n :mw :vw]
930   (sync-scan z3s (make-sliding-stats-mapper 3)))
931
932 | :n | :mw | :vw |
933 |-----+-----+-----|

```

```

934 | 1.0 | 0.857454 | 0.0 |
935 | 2.0 | 0.584954 | 0.148512500000000005 |
936 | 3.0 | 0.6250776666666666 | 0.07908597588033339 |
937 | 4.0 | 0.6190473333333332 | 0.07499115039433346 |
938 | 5.0 | 1.0608326666666668 | 0.2541686787463333 |
939 | 6.0 | 1.05881 | 0.25633817280899995 |
940 | 7.0 | 0.6656836666666668 | 0.9787942981023336 |
941 | 8.0 | 0.04854833333333334 | 0.3215618307563336 |
942 | 9.0 | -0.19849096666666663 | 0.022395237438003604 |
943 | 10.0 | -0.06691730000000007 | 0.01846722403596973 |
944 ... passing the unit test.

```

9 KALMAN FILTER

9.1 BASIC LINEAR ALGEBRA

Go for high performance with CUDA or Intel KML later.

Add the following lines to `project.clj` in the directory that contains this org file:

9.1.1 TODO: FULLY LITERATE: TANGLE PROJECT.CLJ

```

950 [net.mikera/core.matrix "0.62.0"]
951 [net.mikera/vectorz-clj "0.48.0"]
952 [org.clojure/algo.generic "0.1.2"]

```

Smoke test:

```

954 (require '[clojure.core.matrix :as ccm])
955 (ccm/set-current-implementation :vectorz)

956 (ccm/shape
957   (ccm/array [[1 2 3]
958               [1 3 8]
959               [2 7 4]]))

```

3 3

Bits and pieces we will need:

```

962 (ccm/transpose
963   (ccm/array [[1 2 3]
964               [1 3 8]
965               [2 7 4]]))

966 #vectorz/matrix [[1.0,1.0,2.0],
967 [2.0,3.0,7.0],
968 [3.0,8.0,4.0]]

```

`mmul` is multiadic (takes more than two arguments). This is possible because matrix multiplication is associative.

```

971 (let [A (ccm/array [[1 2 3]
972                     [1 3 8]
973                     [2 7 4]])]
974   (ccm/mmul (ccm/transpose A) A (ccm/inverse A)))

975 #vectorz/matrix [[1.0000000000000003,1.0,2.0000000000000004],
976 [2.00000000000000093,3.000000000000001,6.999999999999998],
977 [3.0000000000000006,8.0,3.999999999999999]]

```

9.1.2 DEFN Linspace

```

978
979 (defn linspace
980   "A sequence of $n$ equally spaced points in the doubly closed
981   interval $[a,b]$, that is, inclusive of both ends."
982   [a b n]
983   (let [d (/ (- b a) (dec n))]
984     (map (fn [x] (+ a (* x d))) (range n))))
985
986 (clojure.pprint/pprint (linspace 2 3. 3))

```

(2.0 2.5 3.0)

9.2 DEFN SYMMETRIC PART

```

987
988 (do (defn symmetric-part [M]
989       (ccm/div (ccm/add M (ccm/transpose M)) 2.0))
990     (symmetric-part [[1 2 3]
991                     [1 3 8]
992                     [2 7 4]]))

```

	1.0	1.5	2.5
	1.5	3.0	7.5
	2.5	7.5	4.0

9.3 DEFN ANTI-SYMMETRIC PART

```

994
995 (do (defn anti-symmetric-part [M]
996       (ccm/div (ccm/sub M (ccm/transpose M)) 2.0))
997     (anti-symmetric-part [[1 2 3]
998                           [1 3 8]
999                           [2 7 4]]))

```

	0.0	0.5	0.5
	-0.5	0.0	0.5
	-0.5	-0.5	0.0

```

1001 (let [M [[1 2 3]
1002          [1 3 8]
1003          [2 7 4]]]
1004       (ccm/sub (ccm/add (symmetric-part M)
1005                         (anti-symmetric-part M))
1006               M))

```

	0.0	0.0	0.0
	0.0	0.0	0.0
	0.0	0.0	0.0

9.3.1 DEFN MATRIX ALMOST =

```

1009 (require '[clojure.algo.generic.math-functions :as gmf])

```

The following isn't the best solution: neither relative nor absolute differences are robust. Units in Last Place (ULP) are a better criterion, however, this will unblock us for now.


```

1012 (do (defn matrix-almost=
1013      ([m1 m2 eps]
1014       "Checks for near equality against a given absolute difference."
1015       (mapv (fn [row1 row2]
1016               (mapv (fn [e1 e2] (gmf/approx= e1 e2 eps))
1017                     row1 row2))
1018             m1 m2))
1019      ([m1 m2]
1020       "Checks for near equality against a default absolute difference of 1.0e-9"
1021       (matrix-almost= m1 m2 1.0e-9)))
1022
1023 (let [M [[1 2 3]
1024          [1 3 8]
1025          [2 7 4]]]
1026      (matrix-almost= (ccm/add (symmetric-part M)
1027                               (anti-symmetric-part M))
1028                      M))
1029
1030                                     true true true
1031                                     true true true
1032                                     true true true

```

1030 9.3.2 DEFN SIMILARITY TRANSFORM

```

1031 (defn similarity-transform [A M]
1032   (ccm/mmul A M (ccm/transpose A)))

```

1033 9.3.3 VECTORS, ROW VECTORS, COLUMN VECTORS

1034 The library (like many others) is loose about matrices times vectors.

```

1035 (ccm/mmul
1036   (ccm/matrix [[1 2 3]
1037                [1 3 8]
1038                [2 7 4]])
1039   (ccm/array [22 23 42]))
1040
1041 #vectorz/vector [194.0, 427.0, 373.0]

```

1041 Pedantically, a matrix should only be allowed to left-multiply a column vector, i.e., a 1×3 matrix. The
1042 Clojure library handles this case.

```

1043 (ccm/mmul
1044   (ccm/matrix [[1 2 3]
1045                [1 3 8]
1046                [2 7 4]])
1047   (ccm/array [[22] [23] [42]]))
1048
1049 #vectorz/matrix [[194.0],
1050                  [427.0],
1051                  [373.0]]

```

1051 Non-pedantic multiplication of a vector on the right by a matrix:

```

1052 (ccm/mmul
1053   (ccm/array [22 23 42])
1054   (ccm/matrix [[1 2 3]
1055                [1 3 8]
1056                [2 7 4]]))

```

```
1057 #vectorz/vector [129.0,407.0,418.0]
```

1058 Pedantic multiplication of a row vector on the right by a matrix:

```
1059 (ccm/mmul
1060   (ccm/array [[22 23 42]])
1061   (ccm/matrix [[1 2 3]
1062                [1 3 8]
1063                [2 7 4]]))
1064 #vectorz/matrix [[129.0,407.0,418.0]]
```

1065 9.3.4 SOLVING INSTEAD OF INVERTING

1066 Textbooks will tell you that, if you have $Ax = b$ and you want x , you should compute $A^{-1}b$. Don't do this;
1067 the inverse is numerically risky and almost never needed:

```
1068 (ccm/mmul
1069   (ccm/inverse
1070     (ccm/array [[1 2 3]
1071                 [1 3 8]
1072                 [2 7 4]]))
1073   (ccm/array [22 23 42]))
1074 #vectorz/vector [22.05882352941177,-0.4705882352941142,0.2941176470588234]
```

1075 Instead, use a linear solver. Almost everywhere that you see $A^{-1}b$, visualize `solve(A,b)`. You will get
1076 a more stable answer. Notice the difference in the low-significance digits below. The following is a more
1077 reliable answer:

```
1078 (require '[clojure.core.matrix.linear :as ccml])
1079 (ccml/solve
1080   (ccm/array [[1 2 3]
1081               [1 3 8]
1082               [2 7 4]]))
1083   (ccm/array [22 23 42]))
1084 (ccml/solve
1085   (ccm/matrix [[1 2 3]
1086                [1 3 8]
1087                [2 7 4]]))
1088   (ccm/matrix [22 23 42]))
1089 #vectorz/vector [22.058823529411764,-0.4705882352941176,0.2941176470588236]
1090 (ccm/shape (ccm/matrix [[22] [23] [42]]))
```

```
1091                                     3 1
```

1092 9.3.5 DEFN SOLVE MATRIX

1093 We need `solve` to work on matrices:

```
1094 (defn solve-matrix
1095   "The 'solve' routine in clojure.core.matrix only works on Matrix times Vector.
1096   We need it to work on Matrix times Matrix. The equation to solve is
1097
```

```

1098 Ann * Xnm = Bnm
1099
1100 Think of the right-hand side matrix Bnm as a sequence of columns. Iterate over
1101 its transpose, treating each column as a row, then converting that row to a
1102 vector, to get the transpose of the solution X."
1103 [Ann Bnm]
1104 (ccm/transpose (mapv (partial ccm/solve Ann) (ccm/transpose Bnm))))

1105 (solve-matrix
1106   (ccm/matrix [[1 2 3]
1107                [1 3 8]
1108                [2 7 4]]))
1109   (ccm/matrix [[22] [23] [42]] ))

                                     22.058823529411764
1110                                     -0.4705882352941176
                                     0.2941176470588236

1111 (solve-matrix
1112   (ccm/matrix [[1 2 3]
1113                [1 3 8]
1114                [2 7 4]]))
1115   (ccm/matrix [[22 44]
1116                [23 46]
1117                [42 84]]))

                                     22.058823529411764   44.11764705882353
1118                                     -0.4705882352941176   -0.9411764705882352
                                     0.2941176470588236    0.5882352941176472

```

9.4 DEFN KALMAN UPDATE: GENERAL EXTENDED KALMAN FILTER

Use Clojure's destructuring to write the Kalman filter as a binary function. See <http://vixra.org/abs/1606.0348>

$x_{n,1}$ denotes a vector x with dimension $n \times 1$, that is, a column vector of height n . $P_{n,n}$ denotes a covariance matrix of dimension $n \times n$, and So on.

The math is as follows (notice step 6 has the same form as all earlier statistics calculations in this document):

Letting inputs:

- $x_{n,1}$ be the current, best estimate of the n -dimensional state of a system
- $P_{n,n}$ be the current, best estimate of the $n \times n$ covariance of state $x_{n,1}$
- $z_{m,1}$ be the current, m -dimensional observation
- $H_{m,n}$ be linearized observation model to be inverted: $z_{m,1} = H_{m,n} \cdot x_{n,1}$
- $A_{n,n}$ be linearized dynamics
- $Q_{n,n}$ be process noise (covariance) accounting for uncertainty in $A_{n,n}$
- $R_{m,m}$ be observation noise (covariance) accounting for uncertainty in $z_{m,1}$

and intermediates and outputs:

- $x'_{n,1}$ (intermediate; *update*) be the estimate of the state after enduring one time step of linearized dynamics

- $\mathbf{x}_{n,1}''$ (output; *prediction*) be the estimate of the state after dynamics and after information from the observation $\mathbf{z}_{m,1}$
- $\mathbf{P}_{n,n}'$ (intermediate; *update*) be the current, best estimate of the $n \times n$ covariance of state $\mathbf{x}_{n,1}$ after dynamics
- $\mathbf{P}_{n,n}''$ (output; *prediction*) be the current, best estimate of the $n \times n$ covariance of state $\mathbf{x}_{n,1}$ after dynamics and observation $\mathbf{z}_{m,1}$

The steps are:

1. *Update state estimate*: $\mathbf{x}_{n,1}' = \mathbf{A}_{n,n} \mathbf{x}_{n,1}$
2. *Update state covariance*: $\mathbf{P}_{n,n}' = \mathbf{Q}_{n,n} + (\mathbf{A}_{n,n} \mathbf{P}_{n,n} \mathbf{A}_{n,n}^T)$
3. *Covariance-update scaling matrix*: $\mathbf{D}_{m,m} = \mathbf{R}_{m,m} + (\mathbf{H}_{m,n} \mathbf{P}_{n,n}' \mathbf{H}_{m,n}^T)$
4. *Kalman gain*: $\mathbf{K}_{n,m} = \mathbf{P}_{n,n}' \mathbf{H}_{m,n}^T \mathbf{D}_{m,m}^{-1}$
(a) written as $\mathbf{K}_{n,m}^T = \text{solve}(\mathbf{D}_{m,m}^T, \mathbf{H}_{m,n} \mathbf{P}_{n,n}')$
5. *Innovation: predicted observation residual*: $\mathbf{r}_{m,1} = \mathbf{z}_{m,1} - \mathbf{H}_{m,n} \mathbf{x}_{n,1}'$
6. *State prediction*: $\mathbf{x}_{n,1}'' = \mathbf{x}_{n,1}' + \mathbf{K}_{n,m} \mathbf{r}_{m,1}$
7. *Covariance reduction matrix*: $\mathbf{L}_{n,n} = \mathbf{I}_{n,n} - \mathbf{K}_{n,m} \mathbf{H}_{m,n}$
8. *Covariance prediction*: $\mathbf{P}_{n,n}'' = \mathbf{L}_{n,n} \mathbf{P}_{n,n}'$

```
(defn kalman-update [{:keys [xn1 Pnn]} {:keys [zm1 Hmn Ann Qnn Rmm]})
  (let [x'n1 (ccm/mmul Ann xn1) ; Predict state
        P'nn (ccm/add
              Qnn (similarity-transform Ann Pnn)) ; Predict covariance
        Dmm (ccm/add
              Rmm (similarity-transform Hmn P'nn)) ; Gain precursor
        DTmm (ccm/transpose Dmm) ; Support for "solve"
        HP'Tmn (ccm/mmul Hmn (ccm/transpose P'nn)) ; Support for "solve"
        ; Eqn 3 of http://vixra.org/abs/1606.0328:
        KTmn (solve-matrix DTmm HP'Tmn)
        Knm (ccm/transpose KTmn) ; Kalman gain
        ; innovation = predicted obn residual
        rml (ccm/sub zm1 (ccm/mmul Hmn x'n1))
        x''n1 (ccm/add x'n1 (ccm/mmul Knm rml)) ; final corrected estimate
        n (ccm/dimension-count xn1 0)
        ; new covariance ? catastrophic cancellation ?
        Lnn (ccm/sub (ccm/identity-matrix n)
                     (ccm/mmul Knm Hmn))
        P''nn (ccm/mmul Lnn P'nn)] ; New covariance
    {:xn1 x''n1, :Pnn P''nn}))
```

9.4.1 UNIT TEST

Let the measurement model be a cubic:

```
(defn Hmn-t [t]
  (ccm/matrix [( * t t t) (* t t) t 1 ])))
```

Ground truth state, constant with time in this unit test:

```

1179 (def true-x
1180   (ccm/array [-5 -4 9 -3]))

1181 (require '[clojure.core.matrix.random :as ccmr])

1182 (defn fake [n]
1183   (let [times      (range -2.0 2.0 (/ 2.0 n))
1184         Hmns       (mapv Hmn-t times)
1185         true-zs     (mapv #(ccm/mmul % true-x) Hmns)
1186         zmls        (mapv #(ccm/add
1187                             % (ccm/array
1188                               [(ccmr/rand-gaussian)])))
1189         true-zs     (mapv #(ccm/mmul % true-zs) Hmns)
1190         {:times times, :Hmns Hmns, :true-zs true-zs, :zmls zmls})])

1191 (def test-data (fake 7))

1192   A state cluster is a vector of  $\mathbf{x}$  and  $\mathbf{P}$ :

1193 (def state-cluster-prior
1194   {:xn1 (ccm/array [[0.0] [0.0] [0.0] [0.0]]),
1195    :Pnn (ccm/mul 1000.0 (ccm/identity-matrix 4))})

1196   An obn-cluster is a vector of  $\mathbf{z}$ ,  $\mathbf{H}$ ,  $\mathbf{A}$ ,  $\mathbf{Q}$ , and  $\mathbf{R}$ . Obn is short for observation.

1197 (def obn-clusters
1198   (let [c (count (:times test-data))]
1199     (mapv (fn [zml Hmn Ann Qnn Rmm]
1200             {:zml zml, :Hmn Hmn, :Ann Ann, :Qnn Qnn, :Rmm Rmm})
1201           (:zmls test-data)
1202           (:Hmns test-data)
1203           (repeat c (ccm/identity-matrix 4))
1204           (repeat c (ccm/zero-matrix 4 4))
1205           (repeat c (ccm/identity-matrix 1))
1206           )))

1207 (clojure.pprint/pprint (reduce kalman-update state-cluster-prior obn-clusters))

1208 {:xn1 #vectorz/matrix [[-4.6881351375660065],
1209 [-4.098904857550219],
1210 [7.903839925384],
1211 [-2.657281371213249]],
1212 :Pnn
1213 #vectorz/matrix [[0.03208215055213958,-5.478256737134757E-15,-0.0874691388122202,-8.77076
1214 [-2.3568386825489895E-15,0.03637145347999561,-5.2632377622874316E-14,-0.05541947257604415],
1215 [-0.08746913881223455,-2.570860191397628E-14,0.2822249372573019,-1.1334683192032458E-14],
1216 [4.6455894686658894E-15,-0.05541947257607027,-6.734196533741965E-15,0.15110531309503664]]}

1217   Notice how close the estimate  $\mathbf{x}_{n \times 1}$  is to the ground truth,  $[-5, -4, 9, -3]$  for  $\mathbf{x}$ . A chi-squared test would
1218 be appropriate to complete the verification (TODO).

```

9.5 DEFN MAKE-KALMAN-MAPPER

```

1220 Just as we did before, we can convert a foldable into a mappable transducer and bang on an asynchronous
1221 stream of data. This only needs error handling to be deployable at scale. Not to minimize error handling:
1222 it's a big but separable engineering task.

```

```

1223 (do (defn make-kalman-mapper [{:keys [xn1 Pnn]})
1224     ;; let-over-lambda (LOL); here are the Bayesian priors
1225     (let [estimate-and-covariance (atom {:xn1 xn1, ;; prior-estimate
1226                                           :Pnn Pnn, ;; prior-covariance
1227                                           })]
1228       ;; here is the mapper (mappable)
1229       (fn [{:keys [zm1 Hmn Ann Qnn Rmm]}]
1230         (let [{:keys [xn1 :xn1, Pnn :Pnn]} @estimate-and-covariance]
1231           (let [;; out-dented so we don't go crazy reading it
1232                 x'nl (ccm/mmul Ann xn1) ; Predict state
1233                 P'nn (ccm/add
1234                       Qnn (similarity-transform Ann Pnn)) ; Predict covariance
1235                       Dmm (ccm/add
1236                             Rmm (similarity-transform Hmn P'nn)) ; Gain precursor
1237                       DTmm (ccm/transpose Dmm) ; Support for "solve"
1238                       HP'Tmn (ccm/mmul Hmn (ccm/transpose P'nn)) ; Support for "solve"
1239                       ; Eqn 3 of http://vixra.org/abs/1606.0328
1240                       KTmn (solve-matrix DTmm HP'Tmn)
1241                       Knm (ccm/transpose KTmn) ; Kalman gain
1242                       ; innovation = predicted obn residual
1243                       rml (ccm/sub zm1 (ccm/mmul Hmn x'nl))
1244                       x''nl (ccm/add x'nl (ccm/mmul Knm rml)) ; final corrected estimate
1245                       n (ccm/dimension-count xn1 0)
1246                       ; new covariance ? catastrophic cancellation ?
1247                       Lnn (ccm/sub (ccm/identity-matrix n)
1248                                     (ccm/mmul Knm Hmn))
1249                       P''nn (ccm/mmul Lnn P'nn)]
1250                 (swap! estimate-and-covariance conj
1251                       [:xn1 x''nl]
1252                       [:Pnn P''nn]) ) )
1253         @estimate-and-covariance) ) )
1254
1255     ;; The following line maps over a fixed sequence in memory
1256     #_(clojure.pprint/pprint (last
1257                               (map (make-kalman-mapper state-cluster-prior)
1258                                   obn-clusters)))
1259
1260     #_(async-randomized-scan obn-clusters
1261                             (make-kalman-mapper state-cluster-prior)
1262                             clojure.pprint/pprint)
1263
1264     (let [accumulator (make-sow-reap)]
1265       (async-randomized-scan obn-clusters
1266                             (make-kalman-mapper state-cluster-prior)
1267                             (accumulator ::sow))
1268       (last (accumulator ::reap))) )
1269
1270 '(:xn1 #vectorz/matrix ((-4.688135137565858)
1271 (-4.098904857550392)
1272 (7.903839925383772)
1273 (-2.6572813712131196)) :Pnn #vectorz/matrix ((0.03208215055213749 3.3957212042246E-16 -0.
1274 (4.641605025682005E-16 0.036371453479968716 -7.273641324662128E-16 -0.05541947257606678)
1275 (-0.08746913881223027 -2.2724877535296173E-16 0.2822249372572417 -2.921274333544943E-15)
1276 (1.700029006457271E-16 -0.0554194725760668 -1.6479873021779667E-15 0.15110531309503966)))

```

10 VISUALIZATION SANDBOX

=====

11 VISUALIZATION

>>>>> 893a63642627909f5cdb2aa3ef7083b1e89baa3f :CUSTOM_ID: oz-for-visualization

11.1 CLJ-REFACTOR

```
(list org-babel-default-header-args
      org-babel-default-inline-header-args
      org-babel-default-lob-header-args)

(require 'clj-refactor)

(defun my-clojure-mode-hook ()
  (clj-refactor-mode 1)
  (yas-minor-mode 1) ; for adding require/use/import statements
  ;; This choice of keybinding leaves cider-macroexpand-1 unbound
  (cljr-add-keybindings-with-prefix "C-c C-m"))


(add-hook 'clojure-mode-hook #'my-clojure-mode-hook)
```

Hot-loading seems hopelessly broken from org mode (might work in .clj files).

```
(cljr-add-project-dependency)
```

11.2 INCANTER

```
(use '(incanter core charts pdf))
;;; Create the x and y data:
(def x-data [0.0 1.0 2.0 3.0 4.0 5.0])
(def y-data [4.3 9.0 2.6 3.1 8.1 4.5])
(def xy-line (xy-plot x-data y-data))
#_(view xy-line)
(save-pdf xy-line "incanter-xy-line.pdf")
(save xy-line "incanter-xy-line.png")
```



1304

1305 11.3 OZ

1306 From [https://github.com/metasoarous/oz/blob/master/examples/clojupyter-example.](https://github.com/metasoarous/oz/blob/master/examples/clojupyter-example.ipynb)
 1307 ipynb

```
1308 (require '[clojupyter.misc.helper :as helper])
1309 (helper/add-dependencies '[metasoarous/oz "1.6.0-alpha2"])
1310 (require '[oz.notebook.clojupyter :as oz])
```

1311 11.3.1 DEFN PLAY DATA

```
1312 (do (defn play-data [& names]
1313     (for [n names
1314           i (range 20)]
1315       {:time i :item n
1316        :quantity (+ (Math/pow (* i (count n)) 0.8)
1317                     (rand-int (count n)))})
1318 (def stacked-bar
1319   {:data {:values (play-data "munchkin" "witch"
1320                             "dog" "lion" "tiger" "bear")}
1321    :mark "bar"
1322    :encoding {:x {:field "time"}
1323              :y {:aggregate "sum"
1324                  :field "quantity"
1325                  :type "quantitative"}
1326              :color {:field "item"}}})
```



```

1327 (oz/view! stacked-bar) )

1328 (def spec
1329   { :data { :url "https://gist.githubusercontent.com/metasoarous/4e6f781d353322a44b9cd3e45970
1330     :mark "point"
1331     :encoding
1332     { :x { :field "Horsepower", :type "quantitative" }
1333       :y { :field "Miles_per_Gallon", :type "quantitative" }
1334       :color { :field "Origin", :type "nominal" } } })
1335 (oz/view! spec)

1336 (oz/view!
1337   [ :div
1338     [ :h1 "A little hiccup example" ]
1339     [ :p "Try drinking a glass of water with your head upside down" ]
1340     [ :div { :style { :display "flex" :flex-direction "row" } }
1341       [ :vega-lite spec ]
1342       [ :vega-lite stacked-bar ] ] ] )


```

12 GAUSSIAN PROCESSES

The Extended Kalman Filter above is a generalization of linear regression.

12.1 RECURRENT LINEAR REGRESSION

13 SANDBOX




1348

“

1349

Sede	Max cites	H-index
Chile	257.72	21.39
Leeds	165.77	19.68
Sao Paolo	71.00	11.50
Stockholm	134.19	14.33
Morelia	257.56	17.67



1350

1351 Emacs 26.2 of 2019-04-12, org version: 9.2.2